

SoK: Machine Learning with Confidential Computing

Fan Mo*, Zahra Tarkhani^{†‡}, Hamed Haddadi*

*Imperial College London, [†]University of Cambridge, [‡]Microsoft Research

Abstract—Privacy and security challenges in Machine Learning (ML) have become a critical topic to address, along with ML’s pervasive development and the recent demonstration of large attack surfaces. As a mature system-oriented approach, confidential computing has been increasingly utilized in both academia and industry to improve privacy and security in various ML scenarios. In this paper, we systematize the findings on confidential computing-assisted ML security and privacy techniques for providing *i) confidentiality guarantees* and *ii) integrity assurances*. We further identify key challenges and provide dedicated analyses of the *limitations* in existing *Trusted Execution Environment* (TEE) systems for ML use cases. We discuss prospective works, including grounded privacy definitions, partitioned ML executions, dedicated TEE designs for ML, TEE-aware ML, and ML full pipeline guarantee. These potential solutions can help achieve a much strong TEE-enabled ML for privacy guarantees without introducing computation and system costs.

Index Terms—Machine Learning, Privacy, Integrity, Confidential Computing, Trusted Execution Environment.

I. INTRODUCTION

Machine learning (ML) has been established as the most promising approach to learning patterns from data, and its applications exist across data processing of online surfing, financial data, health care, autonomous cars, and almost every data-driven field around us. The wide applications have also lead ML models being run on a diverse set of devices – from low-end IoT/mobile devices to high-performance cloud/data centers to provide both training and inference services. However, wide applications also open a large attack surface on ML’s security and privacy. Recent research has massively explored the way to attack on ML e.g., model stealing [1], [2], [3], [4], model inversion [5], [6], model/data poisoning [7], [8], [9], data reconstruction [10], [11], membership/attribute inference [12], [13], [14], etc. Vulnerabilities shown by these attacks have increasingly perturbed ML developers and users due to the negative consequences they can cause.

ML faces security and privacy issues mainly due to, firstly, the *complexity of pipelines* that involve many system/software stacks to provide modern features like acceleration. A full ML pipeline covers the collecting of raw data, the complete training and inference phases, the later use of the trained ML model for prediction, and potential re-train and re-use of the ML model. Because data owners, the host of ML computation, the model owners, and result receivers are most likely to be different entities, the pipeline can be segmented across mutually mistrusting individuals and is left with a broad attack surface. Second, the *weak robustness and low interpretability* of ML algorithms deteriorate the security and

privacy problems. The modern ML model is a large-scale nonlinear system that has been progressively adjusted using stochastic gradient descent during training. Until now, we still lack an accurate representation/description of the dynamics of this training progress. Consequently, ML models often have relatively weak robustness, which has been shown in adversarial examples [15] or poisoning attacks [8]), i.e., a small change to the training process could cause large negative impacts that are difficult, or even impossible, to detect.

Confidential Computing has emerged as one promising approach to achieving trustworthy ML to guarantee privacy. Confidential Computing, defined in the Linux Foundation project Confidential Computing Consortium’s whitepapers [16], [17], is the protection of data in use, in addition to data protection in storage and transmission, by utilizing hardware-based Trusted Execution Environments (TEEs). TEE is one of the emerging techniques in enabling isolated and verifiable code execution inside protected memory (sometimes called enclaves or secure world), separated from the host privileged system stacks such as operating system or hypervisor. Both hardware and software changes are made to achieve such isolation on modern processors (i.e., CPUs). Compared to the conventional external secure processor, like the Trusted Platform Module that uses separated hardware from the rest of the board/SoC, TEEs still remain and run on existing processors which then provide much more computational resources. TEEs have become nearly an essential component of modern processors, and various types of TEEs are implemented by e.g., ARM, Intel, AMD, and NVIDIA to achieve Confidential Computing.

The rapid growth of Confidential Computing pushes a line of research work that leverage it for ML [18], [19], [20], [21], [22] as well as commercial products (e.g., Confidential Computing service from Microsoft Azure [23], Amazon Web Services [24], Google Cloud [25], etc.). Although many attempts have been made and several limitations (e.g., still limited TEE resources for ML use) are summarized [26], there is still no systematization of Confidential Computing-assisted mechanisms, challenges, and especially potential solutions. It is still unclear what the critical scientific and engineering obstacles are and how to overcome them for achieving trustworthy ML. In this paper, we present a systematic analysis of Confidential Computing-assisted solutions to alleviate the security and privacy problems in ML.

Organization. The next sections of the paper are organized as: § II describes the pipeline and paradigms of ML and existing

attack surfaces; § III outlines the threat model and key components of Confidential Computing; § IV summarizes challenges and existing solutions of using Confidential Computing for ML; § V further presents the challenges and solutions of guaranteeing ML *integrity* with Confidential Computing; § VI outlines the existing TEE limits in Confidential Computing that hinder its usage in ML; § VII concludes the current research stage and the further work.

II. MACHINE LEARNING

Machine learning (ML) automates the pattern learning from large datasets in the sense that analysts do not need to manually explore data's latent features and their correlations [27], [28]. In this section, we look into ML pipelines, common paradigms, and attack surfaces that exist among them.

A. Pipeline

The ML pipeline codifies and automates the workflow to produce and apply ML algorithms or models. It generally covers most cases when using ML, which consist of multiple sequential steps forming the lifecycle of using ML.

Data preparation. Data are the basis for ML evidently because ML requires tons of inputs to attempt “trial and error” and learn the patterns. In supervised learning [28], all data are labeled, which usually needs extensive efforts of human annotation. In semi-/un-supervised learning [29], [30], [31], parts/all of the data are unlabeled due to the difficulty of labeling. In reinforcement learning [32], data are in the form of sequences of action, observations, and rewards, which are produced during the learning process. No matter in any form, ML is always data-hungry. Data augmentation [33], [34] is necessary to increase the amount of data by adding slightly modified copies or synthetic copies of existing data.

Model training. Most model training aims to acquire a function $f_\theta(x)$, where model parameters $\theta \in \Theta$ the parameter space, capable of mapping an input x to a predicted decision $\hat{y} = f_\theta(x)$, a value near to the true decision y . The form of $f_\theta(\cdot)$ excluding θ is called model hyperparameters, which are determined by chosen model architectures e.g., the type of layers, the number of neurons, etc. Searching of the best θ set usually is achieved by minimizing the loss $\ell(y, f_\theta(x))$ using Stochastic Gradient Descent (SGD) [35], [36] on training sets, i.e., stepped descent by $\theta \leftarrow \theta - \eta \nabla_\theta f_\theta(x)$, where η is the step size multiplier known as the learning rate. The loss is backward passed through the model, which is called backward propagation. After reaching a good performance on the validation set (or the other subset apart from the training set in cross-validation [37], [38]), the model training stops. $f_\theta(\cdot)$'s final model utility/accuracy is reported on the test set.

Model deployment/Inference. By feeding one unseen input data x_{new} into the ML function, one can get the prediction $\hat{y} = f_\theta(x_{\text{new}})$. This is called the inference stage. Inference could happen at both the *model provider* side or the *data owner* side depending on the trust and collaboration form between the providers and the users. In typical cases of ML inference, the model needs to be deployed i) on a server-side device to

support centralized service e.g., ML inference as a service, or ii) on the users' devices for local inference distributively.

Retraining/Updates. ML models require retraining/updating their parameters to adapt to data with unseen features over time. This usually is achieved by fine-tuning the model with newly collected data. Especially, the fine-tuning could focus on the last layers of the model, which is one way to achieve transfer learning [39], [40], as the first layers already generalize well on one type of data (e.g., images in general) [39], [41], [42]. Two examples of such a technique are i) back-propagation on the last layers only [40] and ii) low-shot learning with weights imprinting on the last layer [43]. However, this would depend on whether the pre-trained models have good representatives of the new data. The updated model can be further uploaded to other users or the central server for knowledge sharing [44], [45], [46].

B. Paradigm

Due to the increasing privacy/security concern raised by sharing data, as well as network/computation-resource constraints, the stages of the ML pipeline are typically located to different participating entities for a better trade-off among utility, privacy, and cost. Based on the location of the main ML computation (e.g., training or inference), they can be categorized into: Centralized and Federated machine learning.

Centralized machine learning. In centralized ML, the training or inference computations are conducted at a central place such as a server [47], [48]. For both training and inference, the data have to be collected and stored at/near the center, and then the ML function performs training or inference on the data. In such a case, potential retraining will also happen on the server by collecting new data. Nevertheless, centralized ML raises privacy concerns about others' data, considering that all data need to be collected. Therefore, centralized training/inference is usually chosen for the cases where collected data are public (or non-private), or the data contributors are collaborators or have any form of trust/contracts with the center.

Distributed/federated machine learning. When data owners do not have the willingness to share data, distributed/federated ML is one solution, where the *ML model*, instead of the private data, is sharing among participants [49], [44]. In ML inference, the model owner *distributes* its ML model to end-users, so that users can perform inference locally on their devices. Potentially, they can conduct further fine-tuning on the model for personalization but still for their own usage. In terms of training, the terminology ‘federated learning (FL)’ [44], [45], [46] has been used intensively referring to the paradigm that users send the locally updated model parameters to the server without the need of revealing their data. Unfortunately, recent works have shown that adversaries can still execute attacks to retrieve sensitive information from the model parameters themselves [50], [13], [10].

C. Attack Surface

The attack surface covers many places within the complete pipeline. Here, we categorize existing attack vectors based on

the core underpinning of information security¹ into: i) *confidentiality*, i.e., privacy of data and intellectual property of models, ii) *integrity* of the ML process. The attacks on every vulnerability in the ML pipeline can be viewed in light of confidentiality and/or integrity.

Confidentiality-related attacks. An adversary can be curious-but-honest, i.e., interested in exploring (unauthorized) sensitive information of data and the model, but honestly performing ML training/inference without changing the computation results². These attack vectors usually exist in the pipeline stages of model training, model deployment/inference, or model updates, because the data and models are owned by different participating entities. In centralized ML, without advanced protections, the computation host can directly access incoming data [51], [52].

Similarly, in federated ML, the host who orchestrates all clients' local training can access their updated models and further infer private information about their local data using these models [50], [13], [53]. Several common attacks are i) Data Reconstruction Attacks [10], [50], [54], aiming at reconstructing original input data based on the observed model or its gradients, ii) Attribute Inference Attacks [13], [55], aiming at inferring the value of users' private properties in the training data, and iii) Membership Inference Attacks [53], [12], aiming at learning whether specific data instances are present in the training dataset. On the other side, the *confidentiality of models* can be targeted. A model that has been trained with great effect (e.g., computation power, data cleaning and collection) is considered the intellectual property of its owner, e.g., machine learning as a service (MLaaS) hosted by Amazon Web Services (AWS). In such cases in inference, one possible attack, called model stealing, is to counterfeit the functionality of this model by exploiting black-box access [1], [56], [2], e.g., querying a large number of prediction results.

Integrity-related attacks. An attacker could actively exploit training/inference results to the ML integrity. For example, one well-known adversarial attack [15], [57] adds calibrated noises to one image/audio, which leads to wrong prediction results. However, the input with noisy perturbation is invisible to human beings. In another example, attackers could take full control of wearable BCI (Brain computing interface) devices (e.g., wheelchair, robotic arms, or drone) by compromising input signals to on-device ML models [58]. Although such attacks manipulate input data, it breaks the integrity of ML inference. Previous work has also shown the possibility of model accuracy degradation by only compromising thread scheduling on multi-threaded ML pipeline [59].

Another example is the model poisoning attack in FL that

could lead to Byzantine fault [7]. Such fault can easily occur because the server does not control all clients' local training and cannot verify that their behaviors follow the promised training processes. For example, if one error in values of local updated parameters has been aggregated into the global model, the complete global model will become unusable. Furthermore, one can manipulate the training set e.g., by adding data with calibrated noisy labels [60]. This fools the classifier to have wrong boundaries for some specifically chosen data points. Backdoors can be added in such a way, and the attackers/other users will be able to trigger such backdoors in later use after model deployment [61], [62], [63].

White-box or black-box. Another wide-used way to categorize ML attack surfaces is based on whether one attack requires access to the internal architecture of a ML model [64], [65]. In black-box attacks, model stealing usually starts from the outside without any prior information and aims to learn the model itself. Membership inference attacks on data privacy are usually black-box because it is already much more efficient than white-box attacks as shown in previous research [65], [66]. White-box access does not significantly increase the attack 'advantage' in disclosing membership privacy. The white-box attack includes almost all data reconstruction attacks, some adversarial example attacks, attribute inference attacks, etc. These attacks usually require detailed model parameters/gradients in order to be performed or have reasonable performance. We refer to Papernot *et al.*'s survey [64] for more details about security and privacy vulnerability in ML.

III. CONFIDENTIAL COMPUTING AND TRUSTED EXECUTION ENVIRONMENTS

Confidential Computing [67] is a quickly emerging technology that provides a level of assurance of privacy and integrity when executing codes on data using Trusted Execution Environments (TEE). Nowadays, most Cloud service vendors have started providing Confidential Computing services, e.g., Google Cloud, Microsoft Azure, etc. They usually leverage the TEE solution supplied by processor manufacturers such as ARM, Intel, AMD, and NVIDIA. In this section, we first discuss threat models considered in Confidential Computing, key components, and then software stacks useful for ML application developers.

A. Threat Models

By considering a stronger adversary who has full access to the host privileged system stacks [68], Confidential Computing significantly reduces the trusted computing base (TCB). By default, TEEs assume that the attacker controls the host OS/hypervisor as well as the host applications executing in userspace. She could, for example, access and corrupt the host OS kernel and process resources (e.g., memory, threads, network, files, locks). Hence, the adversary could be the service provider, the device owner itself, malicious third-party software installed on the devices, or a malicious or compromised host OS/hypervisor. Under such an environment, they

¹The full CIA triad includes Confidentiality, Integrity, and Availability, but the attacks specifically aim at ML system's Availability, e.g., Denial-of-Service (DoS) attacks, have not drawn much attention from both industry and academia currently, probably due to the high similarity of availability problems between conventional systems and ML systems.

²Note that if one adversary does not have access to the model and aims to steal ML parameters for other entities or even training data from malicious model queries, they are active attacks

can perform any possible white-/black-box or confidentiality-/integrity-related attacks.

Confidential Computing enables trust between mutually mistrusting components by utilizing TEEs. TEEs provide an isolated environment running in parallel with the host OS (also known as the rich OS), allowing to securely store data and execute arbitrary code on an untrusted device, with reasonably good performance, through protected and measured memory compartments. Thus, one can demand a remote TEE to perform computation on her (sensitive) code and/or data without revealing them to the TEE's host. Few TEEs could also extend the level of Hardware-based isolation to peripherals and persistence storage [69], [70]. The recent commoditization of TEEs both in high-end and low-end mobile devices makes it an ideal candidate to achieve confidentiality and integrity in ML. However, if the code inside an enclave has vulnerabilities, TEEs do not guarantee the protection of the integrity and confidentiality of that computation. Therefore, the attacker could exploit weaknesses in TEEs system [71] and their interfaces/APIs [72], [73].

Moreover, TEEs are vulnerable to various side- and covert-channel attacks [74], [75], physical attacks [76], confused deputy attacks [77], Denial-of-Service (DoS) attacks, etc [78], [79]. These attack vectors are hard to defend against and are not considered within the standard TEE threat models.

B. Key Components of Confidential Computing

Root of trust measurement. Confidential Computing provides hardware-assisted functionalities to establish trust between TEEs and the external computing environment. First, by enabling Root of Trust (RoT) measurement it ensures the integrity of the TEE system [68], [80]. When the processor measures the chain of all critical system software (including boot loader, firmware, OS, hypervisor, and TEE system) before launching in-enclave code, any integrity violation of the initial state could be detected. A compromised code in the chain cannot escape from being measured and therefore from being detected. The measurement needs hardware-assisted cryptographic keys (e.g., hardware unique keys and endorsement keys), which will be attested later to ensure the integrity of the TEE system [68], [80].

Remote trust establishment and attestation Remote attestation [80], [81] allows the user to determine the level of trust/integrity of a remote TEE before transmitting her sensitive data/code. It enables the user to authenticate the hardware, verify the trusted state of the remote TEE, and check whether the intended software is securely running inside the TEE. Attestation could be conducted by a third party besides the user and host [82] or directly between the service provider and enclave owner [83]. For example, the attestation server could be launched by the processor manufacturer/vendor, as in Intel's SGX attestation service [84], which could be different from the host service provider (e.g., a cloud provider)³.

³Nitro Enclave provided by AWS uses its own remote attestation service (e.g., AWS Key Management Service). In such a case, one AWS Nitro Enclave customer provides assurances to their downstream customer.

Trustworthy code execution and compartmentalization.

The core feature provided by Confidential Computing is the hardware-assisted isolation of enclaves/TEEs from the untrusted environment [68]. In general, enclave memory regions are protected and managed through hardware and TEE's system software. Isolation mechanisms for different processor architectures are different (e.g., ARM, Intel, and AMD). For instance, SGX relies on Memory Encryption Engine (MEE) to protect the confidentiality, integrity, and freshness of the CPU-DRAM traffic over enclave memory ranges, while TrustZone relies on separate page tables, hardware privilege layers (e.g., EL3 and Secure EL2/EL1/EL0), and TrustZone address space controller (TZASC)⁴. More importantly, they run enclaves in different privilege and execution modes. For instance, SGX enclaves run in userspace mode, while AMD's SEV and Intel TDX [85] run them on hypervisor mode, and TrustZone secure world runs as a virtual separated core from the host core.

Moreover, depending on the TEE/enclave architecture and threat model, it could support additional peripheral compartmentalization and persistent storage sealing. In Table I we summarize common hardware-assisted TEEs with their providers, supported processors, and features. Among them, TrustZone is developed from a very early period, and it is the most widely-used TEE on mobile/ubiquitous devices. However, it has several constraints including the inflexible protected memory, supporting only one single secure world, and restricted third-party application adaption. To overcome these issues, Recently ARM proposed Confidential Compute Architecture (CCA) which aim to run more flexible enclaves (called realm) in parallel with TrustZone. Therefore, modern TEEs support multi-enclaves and more flexible secure memory sizes with different TCBs. TEE systems which include an OS-like software stack usually have a large TCB size. On the other side, a larger TCB size can usually provide higher usability to developers and users. For instance, when TEEs provide such an OS, they statically link applications into the in-enclave *kernel*; consequently, unmodified applications could be deployed into enclaves.

C. Partitioning Frameworks for Developers

With the low-level hardware and system stacks only, it is still hard for application developers to utilize TEEs. Hence, depending on the use-cases, various software vendors provide deployment and partitioning frameworks for running a diverse set of workloads within enclaves. For example, more recently Amazon Nitro enclaves and Enarx support running the same binary within different types of TEEs through in-enclave WebAssembly (WASM) support, inspired by Haven [99] that ports Drawbridge [100] (a Windows library OS) inside an SGX enclave. Similarly, Graphene-SGX [101] ports the Linux-based Graphene library OS in an enclave. Scone [102] tries

⁴There are software/language-based TEEs and compartmentalization techniques without the needs of hardware-based guarantees (e.g., seL4), but we do not consider them here as they usually have lower security assurance and are not the mainstream in Confidential Computing.

TABLE I: Summary of dominant hardware-assisted trusted execution environments

Main Stream Hardware TEEs	ISA & Year	Supported Processors	Number of Isolation	TCB Size (excluding CPU/SoC package) ^α	Secure Memory Size ^δ	Attestation	Application	Application Unchanged	Protections
Software Guard Extensions (SGX)	Intel (2013)	6th Intel CPU +	multi-enclaves	small. BIOS or firmware	up to 128MB	remote Intel service	desktop-level	○	confidentiality Integrity
TrustZone (TZ)	Arm (2005)	ARMv6-A + / ARMv8-M +	single secure world	small. firmware, TZ-kernel	typically up to 16 ~ 64MiB	TPM-based [86]	mobile-level desktop-level	○	confidentiality integrity
Secure Encrypted Virtualization (SEV) ^β	AMD (2016)	AMD EPYC +	multi-VMs	large. VM's OS, firmware	up to available system RAM	secure processor	enterprise-level	●	confidentiality
PMP-based TEE ^β	RISC-V (2017)	Xilinx Artix-7 SiFive E31, U54	multi-enclaves	changeable. Runtime, firmware	changeable	-	mobile-level desktop-level	○	confidentiality integrity
Nitro Enclaves	AWS ^γ (2020)	Nitro Cards & Security Chip (EC2)	multi-enclaves	changeable. VMs, OS/Hypervisor	up to available instance RAM	remote AWS service	desktop-level enterprise-level	●	confidentiality integrity
EdgeLock Secure Enclave	NXP (2021)	i.MX 8ULP(-CS) i.MX 9	single secure world	small. Firmware	changeable typically small	remote Azure Sphere	IoT-level mobile-level	○	confidentiality integrity
Confidential Compute Architecture (CCA)	Arm (2022)	Arm v9-A	multi-realms	changeable. OS kernel, firmware	up to available system RAM	remote	mobile-level desktop-level	●	confidentiality integrity

^α Full TCB size is implementation-specific, and the reported TCB size here is the best case supported by hardware (bare-metal potential TCB). ^δ Most difficult for fitting ML based on prior works.

^β Strictly speaking, PMP (physical memory protection) is not a TEE but a hardware-based memory isolation feature that enables several frameworks to architect different TEEs on RISC-V.

^γ Nitro adds new infrastructure but still builds atop existing CPUs so strictly not an ISA. Application Unchanged: ○ Major changes needed; ● Light changes needed; ● No change needed

TABLE II: Open source TEE frameworks

Open Source TEE Framework	Leading	Supported TEEs	On Top Of	Main Functions	Attestation	Languages	ML Examples
Intel SGX SDK [87]	Intel	SGX	firmware & driver	API and libraris for SGX app.	●	C/C++	●
OP-TEE [88]	Linaro	TrustZone	firmware & driver	API and libraris for TZ app.	○	C/C++	●
MultiZone [89]	HEX-five	RISC-V	firmware & driver	API and libraris for RISC-V Enclave app.	●	C/C++	○
Keystone [90]	UC Berkeley	RISC-V	firmware & driver	API and libraris for RISC-V Enclave app.	●	C/C++	●
Open Enclave SDK [91]	Microsoft	SGX, TrustZone	SGX SDK & OP-TEE	generalized API for TEE app.	●	C/C++	●
Asylo [92]	Google	SGX	SGX SDK	generalized API for TEE app.	●	C/C++	●
Teaclave [93]	Apache	SGX	SGX SDK	API for Rust programming	●	Rust	●
Gramine [94]	Invisible Things Lab / Intel	SGX	SGX SDK	lightweight library OS to host app.	●	Native executable*	●
Occlum [95]	Tsinghua / Ant Group	SGX	SGX SDK	lightweight library OS to host app.	●	Native executable*	●
SCONE [96]	Scontain	SGX	container & driver	TEE-based container image generation	●	Docker image*	●
Enarx [97]	Red Hat	SGX, SEV, CCA (upcoming)	firmware & driver	WebAssembly sandbox	●	Wasm binary*	○
Veraencz [98]	Arm	SGX, TrustZone, Nitro Enclaves, CCA (upcoming)	SGX SDK & OP-TEE Nitro Enclaves SDK	proxy attestation, policy WebAssembly sandbox	●	Wasm binary*	●

Attestation: ○ Not available; ● Self-configuration required; ● Work off-the-shelf * Most modern languages supported by compiling as a target executable/binary.

ML examples: ○ Not exist; ● Exists However, One non-exist example can be developed using an existing ML framework with corresponding languages.

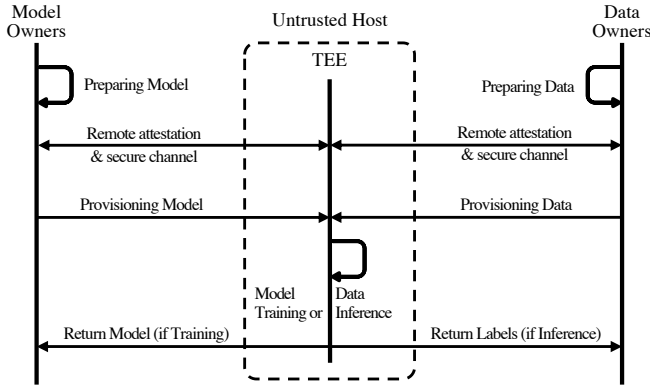


Fig. 1: Overview flow of Confidential Computing which utilizes the untrusted host’s Trusted Execution Environment to protect the model and data in machine learning.

to reduce TCB by porting `musl libc` and a portion of Linux Kernel Library (LKL) [103]. TrustShadow [69] also protects unmodified applications from the host OS by trapping applications exceptions and system calls inside TrustZone, passing the calls to the OS, and verifying the output after parameter marshaling.

As another approach, various TEE partitioning frameworks split applications into trusted and untrusted components, e.g., Intel SGX SDK [104], Microsoft Open Enclave [105], Google’s Asylo [106], OP-TEE [88], and Keystone [107]. There are also language-specific partitioning frameworks such as Civet [108] for porting Java classes into SGX enclaves, Trusted Language Runtime [109] for running portions of C# applications inside TrustZone, and Glamdring [110], a compiler for partitioning applications into SGX enclaves via code annotation. Table II summarises primary commercial TEE frameworks. However, not all of them are feature-rich enough to support ML use cases. Particularly, mobile vendors (e.g., TEE of Qualcomm [111], Trustonic [112], or Huawei [113]) also only allow for limited TEE operations to avoid security risks inside TrustZone secure world [77]. Besides, such security concerns link to the less privileged execution of enclaves as well as ease of programming.

IV. CONFIDENTIAL COMPUTING SOLUTION FOR MACHINE LEARNING

A. Overview

In Confidential Computing-assisted ML, data/model owners need to secretly provision their data/model to the untrusted host’s TEE (see Figure 1). Specifically, after preparation of the model and/or data, the owners first perform remote attestation to assure the integrity of the remote TEE, and then, establish secure communication channels to the TEE. Afterward, data/models are provisioned to the TEE, where model training or inference will be performed. The final results will be returned, i.e., a trained model will be transmitted out in training, or the data label will be returned back to the user in inference. Note that in practice the host, model owners, data

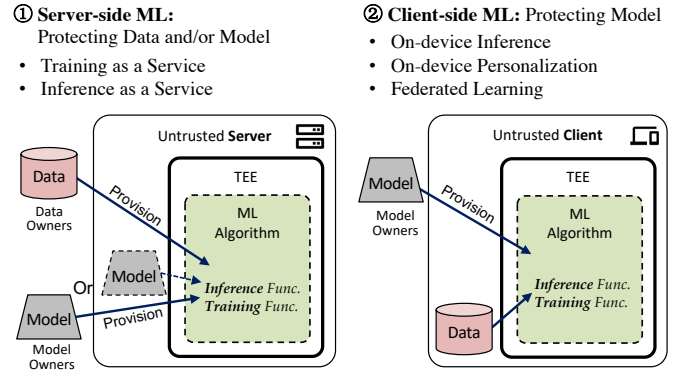


Fig. 2: Server-side ML (Left) and client-side ML (Right) protection using Trusted Execution Environments.

owners, and result receivers could be distinct entities, or some may not exist depending on use cases. We further classify these cases according to the nature of the host: Server or Client (see Figure 2).

Server-side ML protection. A server-like host aims to provide ML service to its customers. Based on the specific ML functionalities, the host can provide i) Inference as a Service (e.g., [114], [21]) or ii) Training as a Service (e.g., [20], [115]). In both cases, schemes should protect *data privacy*. That is, data have to be secretly provisioned to the untrusted server’s TEE no matter for inference or training. The private information contained in the model⁵ or its *intellectual property* also requires protection if it is owned by another entity (e.g., external model owners/providers), not the server. Then the model will be secretly provisioned into the TEE as well. However, note that the ownership of the ML model can belong to not only external owners but also to the untrusted server itself. The latter case will only exist for ML inference, as performing inference does not require changing the model parameters. In contrast with it, training leads to updates of the model parameters on data which consequently leaks private information about the data (i.e., *data privacy*).

Client-side ML protection. A client-like host usually acts as the downstream user of a server, and it performs ML based on the ML framework/algorithm and models provided by the server. Common use cases are i) on-device inference (e.g., [116]), the client conducts predictions on its own data with another individual’s model, ii) on-device personalization (e.g., [117]), the client personalizes the model based on its data for later inference, and iii) federated learning (e.g., [22]), the client trains the model to further contribute a global model owned by another party (e.g., server). In all these cases, the client owns the data; thus, this client’s TEE does not need to hide these local data from itself. However, the model can be considered confidential as it could be *intellectual property*, e.g., the model costs its owner or the server significant effort to

⁵Note that we do not use ‘model privacy’ to avoid confusion, as the private information from ML model is basically about the data used to train it, i.e., in a form of data privacy. The effort of building a model from these data is more proper to be called intellectual property.

train. More importantly, e.g., the model itself can leak private information about its previous trainers (server or other clients). The TEE of one client then provides a trusted environment to respect other clients' data privacy.

B. Key Challenges of Adapting Confidential Computing

The key challenges mostly lie in the feasibility and effectiveness of utilizing TEEs for various ML services due to TEEs' limitations.

Constrained execution environments. TEEs have constrained resources for reducing the size of TCB. As given in Table I, the commonly used hardware TEE, SGX, provides 128MB secure memory⁶, while TrustZone-based TEEs, such as OPTEE or trusty, by default provide up to 16~64MiB secure memory. Depending on the TEEs software stack and architecture, the size and flexibility of protected memory vary. However, such memory sizes are still highly limited compared to the memory consumed by current ML algorithms, which can reach hundreds of MBs or GB. Although techniques using page swapping can increase available memory, it leads to significant overhead (e.g., 100~1000 \times) when running ML [21], [119].

In addition, processor capabilities can be limited when running in the TEE mode. For example, privilege instructions are not allowed inside SGX enclaves, which makes supporting widely used functionalities (e.g., filesystem or multi-threading) difficult or even impossible. For instance, the SGX multi-threading model relies on untrusted OS for thread scheduling which could lead to synchronization-based attacks (e.g., AsyncShock [120]). Similarly, making the secure world over-privileged in TrustZone architecture, caused mobile vendors to only provide limited TEE services and much-restricted TEE usage. Indeed, increasing computational resources is always preferable for ML developers but it can reduce the reliability and practicality of TEE-based solutions.

SDK supports. Modern ML framework usually requires numerous libraries and cross-compilation to support high-performant data loading and computations (e.g., matrix multiplication). However, most TEEs provide basic low-level SDKs. Even though open source TEE frameworks have been developed to support sandbox/container (see Table II), it is still hard to port all ML dependencies and necessary libraries into TEEs. For example, WebAssembly supports compiling binaries from many modern languages, e.g., Rust, which could have available ML libraries, such as autograd [121], but they have much-limited functionality compared to Tensorflow or PyTorch [122]. OS-based containers can provide richer libraries, but still, the lightweight version OS cannot have a rich environment same as the normal OSs like Linux. Porting a normal OS to TEEs is also infeasible (and overkill in many use cases) because it either exceeds the TEEs' size limits or leads to a huge TCB which destroys TEEs' security benefits.

Privacy protection effectiveness. By deploying ML inside TEEs (for training/inference), the untrusted host itself cannot

access the ML computations anymore. However, running inside TEEs does not automatically disable all potential privacy leakages. The other parts of the ML pipeline could still leak private information through various attack vectors. For instance, the produced result, e.g., predicted data labels or trained models, usually need to be transmitted out of the TEE through potentially untrusted channels (e.g., network, IPC, or file system) which could be used to exploit private information [10], [13], [123], [124], [125]. Thus, one particular strategy could have inconsistent protection effectiveness for other ML scenarios and different types of defined privacy [126], [127]. As one example, while hiding model parameters (along with gradient updates and activations) inside TEEs can defend against DRAs, it has very low efficiency for MIAs as most membership information can leak from the model's outputs, i.e., prediction results transmitted out of the TEE [123], [65], [128]. Therefore, for one specific TEE-assisted use case, the target privacy should be defined clearly, and the achieved protection efficiency also requires explicit analyses.

C. Existing Solutions of Confidential ML

Previous research has been dedicated to achieving ML with TEEs, and some aim to overcome the above challenges. We summarize previous research in Table III.

Complete ML training/inference inside TEEs. The most straightforward approach is to deploy a complete ML training/inference process inside TEEs. In such a case, the maximum capability of the ML task is strictly limited to the TEE's space and computing constraints. Then, the strategy is to use memory space efficiently – trading off the total number of the model's layers and the number of neurons – and to maximize the efficiency of every "bit" of TEEs' secure memory for ML computation. Specifically, we list several practical approaches as follows. i) Conducting inference instead of training. Training consumes much more computational resources because of backward propagation (e.g., memory used to save model gradients and intermediate activations). ii) Choosing a small batch size. A large batch size leads to large memory usage because every sample in this batch produces its own activations for all model layers. iii) Balancing the feature extractor (e.g., convolutional layers) and the classifier (e.g., fully connected layers). A properly designed feature extractor can decrease the feature dimension but still capture key features, enabling a compact classifier to achieve good performance.

Partitioned execution. To avoid exceeding the maximum secure computational resource and memory swapping, one approach is efficient compartmentalization (or partitioned execution) to actively optimize the memory usage in ML. Figure 3 shows two types of partitioned execution: i) Layer-based partition [128], [117], [22], [139], which works for models with layer architecture in general, and ii) Feature map-based partition [134], [135], which specifically aims to convolutional layers because of their high memory cost. Among them, whether to keep the first part or the last part inside TEEs (cases ① and ②) depends on the security/privacy protection


⁶The SGXv2 [118] extensions allow additional flexibility in runtime management of enclave resources (e.g., adding memory to an enclave after the enclave is built and running) and multithreaded execution within an enclave.

TABLE III: Previous research that uses Trusted Execution Environments to guarantee the confidentiality of machine learning.

Work & Year	Confid. Computing		Machine Learning					Advanced Features
	TEE Type	SDK	DL Library	Train-ing	Locat-ion	Protect Aim	Overhead (ex-ecution time) ^α	
Obliv. MP'16 [18]	SGX	Intel SGX SDK	fast CNN	●	S	D	1 – 3% on Native	Data-obliviousness
Chiron'18 [20]	SGX	Intel SGX SDK	Theano	●	S	D	4 – 20% on Native	Multi-enclaves, Ryoan VM
PRIVADO'18 [114]	SGX	Intel SGX SDK	ONNX	○	S	D	~ 17.18% on Native	Compiler, data-obliv.
Myelin'18 [115]	SGX	Intel SGX SDK	TVM	●	S	D	4.7 – 9.2% on Native	DP for data obliviousness
Slalom'18 [19]	SGX	Intel SGX SDK	Eigen	○	S	D	SU. 4 – 20× on SGX	Freivald's algorithm + GPU
Graviton'18 [129]	GPU-TEE	CUDA RT driver	Caffe	●	S	D	17 – 33% on GPUs	secure CUDA for TEE on GPUs
Occlumency'19 [21]	SGX	Intel SGX SDK	Caffe	○	S	D	~ 72% on Native	On-demand loading, Channel partition
TensorSCONE[119]	SGX	Intel SGX SDK	TF	●	S	D	~ 3× on Native	Docker supported, Compiler
YerbaBuena'19 [130]	SGX	Intel SGX SDK	Darknet	○	S	D	up to 7.5% on Native	Layer-wise partition, Privacy measure
OMG'20 [116]	TrustZone	SANCTUARY	TFLM	○	C	M	2.11× on Native	-
DarkneTZ'20 [117]	TrustZone	OP-TEE	Darknet	●	C	M	3 – 10% on Native	Layer-wise partition, Privacy measure
TrustFL'20[131]	SGX	Intel SGX SDK	TF	●	C	TI	~ 2× on Native	GPU-outsourcing, Random sampling
Telekine'20 [132]	GPU-TEE	ROCm & CUDA	MXNet	●	S	D	10 – 41% on GPUs	Timing attacks, data-obliv.
MLCapsule'21 [133]	SGX	Intel SGX SDK	Eigen	○	C	M	55 – 116% on Native	-
Trusted-NN'21 [134]	TrustZone	OP-TEE	-	○	C	M	22.8% on Native	Weights & Feature-map partition
Mem-Eff.'21 [135]	SGX	Azure CC (VM)	Darknet	○	S	D	9 – 100% on Native	Channel & Y-plane partition
PPFL'21 [22]	SGX + TZ	OE + OP-TEE	Darknet	●	C + S	D	~ 15% on Native	Layerwise training, Privacy measure
Goten'21 [136]	SGX	Intel SGX SDK	Eigen	●	S	D	SU. 6.84× on SGX	GPU-outsourcing, Non-colluding servers
Citadel'21[137]	SGX	SCONE	TF	●	S	M + D	9 – 73% on Native	Multi enclaves for training
Fair. Audit'22 [138]	SGX	SGX-LKL-OE	PyTorch	○	S	M + D	~ 30% on Native	Fairness audit, Modler + regulator

^α The reported overhead is given to show general performance and *cannot* be cross-compared due to measuring under different experimental settings. SU. refers to speed-up, instead of overhead. Training: ○ Not available; ● Available Location: S = Server; C = Client Protect Aim: D = Data privacy; M = Model intellectual property; TI = Training integrity

Layer-based Partition



Feature map-based Partition




Fig. 3: Partition on the ML process (shown with model architectures) in order to protect it or a part of it inside TEEs.

goal. Rolling on layers inside TEEs (case ③) naturally works for ML inference, because one input propagates forward throughout layers and never goes back; that is, it works in the way that the next layers are loaded into a TEE until one layer's computation executes finished in the TEE. Rolling on feature

maps (cases ④ and ⑤) reduce the realtime memory usage by applying GEMM (i.e., General matrix multiply) and `img2col` (i.e., Image to column format transform) functions on divided sections and channels of feature maps. However, this disables parallelization features and can increase computation time. Furthermore, since training involves backward propagation where more resources are required, partitioned execution is more necessary than inference. Partitioned execution for training could lead to higher overhead than inference in ①, ②, and ③ (increased from 10% to 20% approximately [117], [22]), because a higher volume of buffer will be transmitted between the TEE and normal OS. Note that in order to enjoy higher flexibility to support generic ML models, the partition should better not change the model architecture (unlike the model partition in the design of AlexNet [140]),

TEE-assisted accelerators. In addition to optimizing the limited TEE-protected resources, one can also extend the computation power and resources to accelerators (e.g., GPUs, NPUs, TPUs, and FPGAs) [129], [136], [141], [142], [143]. However, most accelerators are considered untrusted and do not support TEE functionalities. To address this issue, one way is to use CPU-TEE (e.g., SGX) to coordinate the computation delegation and to establish a secret sharing protocol between the CPU-TEE and untrusted accelerators. In such a way, accelerators are able to perform executions on a “nonsensitive version” of the TEE computation, e.g., masked by one-time pads [19], [136]. This can also accelerate the ML's linear

computation i.e., matrix multiplication. There are also TEE-assisted record-and-replay based solutions, which are more suitable for resource-constrained and mobile GPUs. As an example, CODY [143] allows the full computation to be done inside GPU and uses TEEs like TrustZone for secure logging and inference monitoring. However, due to the complexity of heterogeneous SoC architectures, software-only secure channels between TEEs and accelerators without proper hardware-based root-of-trust weaken TEE security guarantees.

Another approach is to extend accelerators hardware, systems stack (e.g., the driver), and API support for enabling accelerator TEEs. Previous works [129], [132] achieve GPU-TEE by relatively small hardware modification (unmodified GPU/CPU core but modified peripherals like PCI) and the necessary system software (e.g., runtime, hypervisor, or drivers, etc.) These solutions could add up to 41% overhead on native GPUs. Similarly, SGX-FPGA [142] builds a secure hardware isolation path between CPU and FPGA to protect the sensitive data in both components and in interactions. It extends the capability of the original SGX to a counterpart FPGA enclave by leveraging the FPGA’s physical unclonable function (PUF) to achieve the hardware RoT on the heterogeneous SoC. This approach is also integrated into some recent NVIDIA GPUs, such as H100 Tensor Core GPU Architecture [144]. Therefore, despite the current limitations, Confidential Computing would be supported on future accelerators while ML frameworks/solutions are not prepared yet.

Attack-based privacy measurement. The TEE protection needs measurements to show its effectiveness. Currently, the main approach for measuring TEEs security is through red-teaming and performing various attack vectors directly. Whether to and which attack to perform mostly depends on the threat model defined for the TEE which could be inadequate for different ML use cases. As one example, regarding the original input protection with TEEs, one performs attacks to reconstruct input data [128] while there is no need to conduct membership inference attacks. By contrast, TEE protection on membership information does not take care of other private information and only be measured with membership inference attacks [117]. However, due to the poor definition of various types of privacy, the attack-based analysis lacks a theoretical foundation on privacy leakage which sequentially makes the TEE-based privacy guarantee less trustworthy.

V. CONFIDENTIAL COMPUTING HELPING MACHINE LEARNING INTEGRITY

Intuitively, other than privacy guarantee, deploying ML training/inference process into TEEs could avoid malicious modifications to this process and therefore ensure integrity. However, in practice, as the other parts of the ML pipeline still can be breached, protecting only the training/inference stage does not always guarantee integrity. In this section, we present the key challenges in enabling ML integrity using TEEs and the solutions that exist or can be adapted.

A. Key Challenges in Enabling ML Integrity

Large attack surface. There is a large attack surface to breach integrity; however, one cannot expect to deploy the complete ML lifecycle inside TEEs to achieve an integrity guarantee due to the TEEs’ resource constraints. These limitations for achieving a proper integrity guarantee are even more severe compared with the issue we met when guaranteeing confidentiality. Specifically, while a curious-but honest (i.e., confidentiality-related) adversary has the goal to explore specific private information, the integrity-related adversary can be more diverse – it aims at not only actively changing the ML process but also even breaking the process by covertly changing one bit in updated weights/gradients. Indeed, to disclose confidential information, adversaries usually explore the target data/model and contained information in the model may also gradually decrease along with more aggregated information in training. However, to breach integrity, one just needs to change one step in the training process (even in one bit) among participants (e.g., Byzantine attacks [7] or synchronization attacks [59]).

Uncontrolled input/output space. While deploying ML training/inference inside TEEs avoids unauthorized direct changes in the produced result/models, the upstream pipeline, i.e., data preparation or input space, is uncontrolled. Editing the inputs and their labels can affect the later ML process (e.g., called “dirty” inputs/labels), which is also the practice of performing some attacks, e.g., poisoning attacks or adversarial example attacks. Adversarial examples [15], [57] to add calibrated noises to the input can seriously perturb the integrity of the input space due to the hardness of detecting the noisy perturbation. In addition to that, the downstream pipeline can be edited maliciously. For example, after the ML algorithm inside the TEE makes a prediction and transmits this result out, a server-side adversary fakes the result, which is also hard to detect on one result receiver side.

Low interpretability of ML. One major reason for the difficulty of detecting integrity perturbations is the low interpretability of current ML/DL [145]. Specifically, current practitioners train ML models using SGD, i.e., weights of neurons are updated iteratively to find the optimal solution. There is still no comprehensive theoretical framework to fully interpret the internal steps of ML training. Therefore, one cannot determine whether the weight value of a neuron is uncompromised or not without *re-running* the ML process using the same setting. With such low interpretability, malicious changes such as the adversarial perturbation in inputs and the changes in model training are hard to trace. Algorithm-based approaches have been proposed to alleviate this situation e.g., by removing outliers that are considered as unreasonable bits [7], by further training the model on adversarial inputs [146], or by cryptography protocols [147]. However, utilizing a system-based isolation technique (i.e., TEEs) to deal with the interpretability issue is not straightforward. One may still need to find a way to identify and assure key elements of the ML pipeline are trustworthy using TEEs.

B. Existing/Adoptable Solutions

Assurance mechanisms. To assure the integrity of ML training/inference, one way is to rerun the complete or a part of the process inside TEEs. Foreexample, the heaviest computation in ML is the matrix multiplication which is preferable to be outsourced to distrusted GPUs for acceleration. To verify whether this matrix multiplication is honestly performed, one can use Freivald’s algorithm which utilizes randomization to reduce time complexity from $\mathcal{O}(n^{2.373})$ (best-known matrix multiplication algorithm) to $\mathcal{O}(kn^2)$ with a probability of failure less than 2^{-k} [148], [19]. Besides, based on the nature of iterative training of ML, one can also sample and verify only a fraction of the training process in TEEs to reduce the overhead [131], [149]. This will require backups for every several training iterations to reach a checkpoint faster, where Merkle hash tree-based method [150] can be used to reduce the storage overhead. This achieves less than 1% soundness errors even for the adversaries who have honestly completed 90% of training rounds [131]. There could be more efficient ways of sampling or verification based on watermarking [151]. How such an assurance guarantees the integrity of other stages in the ML pipeline deserves more investigation.

Input/output space control. The integrity of input/output in the ML pipeline needs to be assured considering that the input/output space is far more uncontrollable than the training/inference process. One way that is generally used for data assurance and can also be adopted in ML input control is to require the digital signature for the data generation. This can be done by hashing the generated sensor data (inside TEEs for example) to produce hash values. After that, sensor data are authenticated using these hash values [152], [153], [154]. At a later stage, the digital signature would also work in a similar way for remote attestation when outputting prediction results, considering that the result receiver trusts the TEE’s behavior. Nevertheless, it becomes harder especially for supervised learning because the data generation usually involves human annotations that are out of control of the digital signature.

Even in semi-supervised/unsupervised learning, it is possible to generate adversarial data without breaching sensors/-TEEs, e.g., by changing physical surroundings or doing abnormal behaviors that the ML does not expect to learn. Therefore, the exploration of how to assure the integrity of the input space using TEEs is still a research question. One example can be adopting detection mechanisms (e.g., [155], [156], [157] to identify fake/dirty inputs inside TEEs, which deserves further investigations. Previous work also introduced system-wide TEE-assisted information flow control, particularly over peripherals and OS services to detect and mitigate such unauthorized IO-based attack vectors [158]. However, these techniques require fundamental changes in system software including commodity OSs.

VI. LIMITATIONS OF EXISTING TEE SYSTEMS

Despite the benefits of TEEs/enclaves, the right abstractions for securely and efficiently utilizing them are still not clear.

Particularly, when combining various TEE solutions (e.g., for multi-platform use cases) or integrating with existing systems. Below we further detail the limitations of existing TEE systems that obstruct their use in ML.

A. Widening the Attack Surface

New features lead to new attack vectors. TEEs require rich functionalities and performance features to support running modern ML programs. However, merely adding hardware features without a comprehensive security model and mechanism with existing system software and privilege boundaries (e.g., sharing data or resources) can even increase the system’s complexity and open new attack vectors. For example, ARM TrustZone introduced hardware support for secure and non-secure worlds by adding several privilege levels in each world. However, the architecture limitations, such as insecure sharing mechanism and fixed/one-way security model, combined with insecure software designs lead to new attack vectors, particularly when sharing resources or exchanging data/control between different privilege layers [77], [78]. Figure 4 demonstrates some of these attack vectors in a TrustZone-assisted environment. Previous attacks, such as Boomerang [77] and HPE (Horizontal Privilege Escalation) [159], show how the semantic gap between heterogeneous compartments (e.g., userspace processes and TrustZone TAs (Trusted Apps)) leads to severe privilege escalation threats. Boomerang exploits a confused deputy vulnerability inside a TA via the shared memory between the two worlds and takes advantage of secure world privilege to make the host kernel memory accessible. Similarly, HPE attacks demonstrate that an adversary process could compromise another userspace process through a shared TA or insecure privilege management between two TAs.

Huge TCB caused by in-enclave LibOS. Moreover, in-enclave LibOS approaches for complex applications result in a huge TCB through porting all dependencies inside the enclave. These designs force developers to run large code in a single enclave, resulting in inefficient and over-privileged enclaves. This leads to wide range of new attacks due to exploiting in-enclave vulnerabilities [120], [72], insecure interactions with outside [160], [161], [72], or misusing enclaves to escalate privileges [77], [159], [162]. Currently, there is no systematic way to neither detect nor protect against these threats.

Vulnerable interactions with outside. None of the existing partitioning frameworks fully consider the complex attack surface originating from insecure interactions between the host OS, userspace processes, and enclaves. As a simple example, Vicarte *et al.* [59] show how asynchronous poisoning attacks on TEE systems lead to changing the accuracy and integrity violation of in-enclave ML models. Civet uses dynamic taint-tracking to control the flow of objects on enclave interfaces, but can not help for proper privilege separation and against more complex attacks (e.g., horizontal privilege escalation (HPE) attacks [159]) and other languages. Moreover, they do not consider in-address space compartmentalization the key issue of handling over-privileged enclaves. Sirius [158] shows

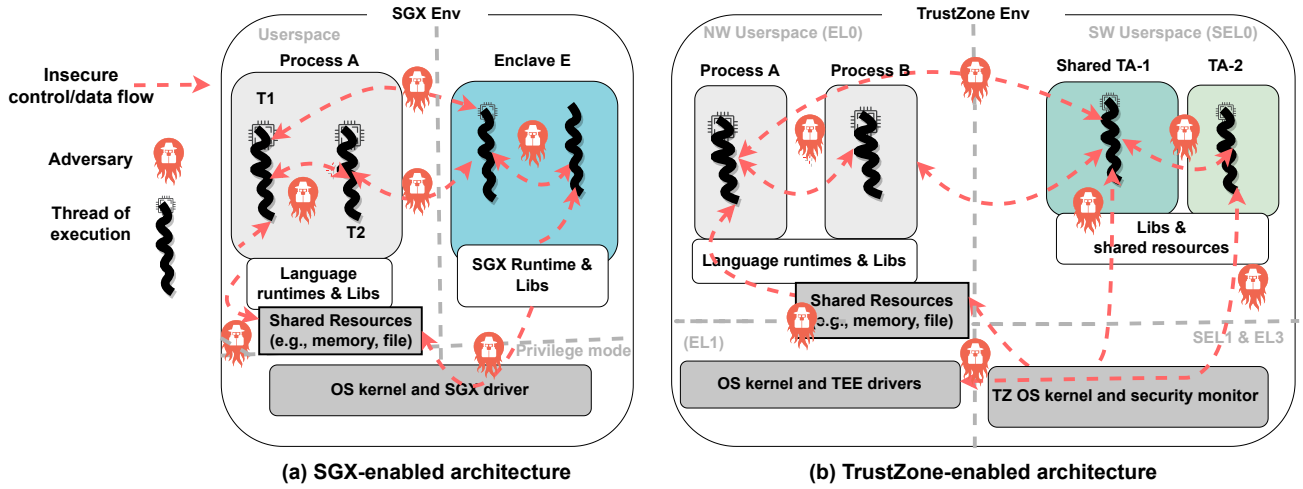


Fig. 4: Simplified attack vectors within and across trust boundaries in TrustZone- and SGX-based TEE systems.

targeting these attacks require fundamental changes to underlying OS and TEE systems. Note that developers sacrifice performance and rich functionality for the security benefits of TEEs; which could be a huge cost without actual benefit considering the insecure designs of existing TEE systems.

ML inside TEEs but still being attacked. Existing vulnerabilities of TEEs still leave the deployment of ML under some forms of attacks. For example, by observing the access pattern to TEEs, previous work in Privado [114] can disclose the target class information i.e., classifying encrypted inputs with high accuracy (97% and 71% for MNIST and CIFAR10 respectively). Also, side-channel attacks on TEEs like timing, power, and Electromagnetic analysis could compromise the ML model privacy on some levels, e.g., reconstructing the model architecture and even model parameters [163], [164], [165]. Some of such attacks can be out of scope because our threat model does not include side-channel attacks. However, this still reflects the current limitations of the existing TEE system which deserve further investigations in future design (e.g., efficient oblivious operations for ML [18], [166]).

B. Programmability & Performance

The quality of ML solutions highly depends on the proper processing of an enormous amount of data. However, almost all TEE systems are not designed considering this essential requirement. Many TEEs lack minimal hardware resources for efficient ML programming on modern CPUs as we explained earlier (Section IV). Currently, resolving these limitations causes significant additional performance overhead. As an example, MPTEE [167] proposes a mechanism for supporting more flexible memory management and dynamic permission enforcement inside enclave memory. It uses MPX's three bound registers to offer the six common memory permissions (i.e., RWX, RW, RX, R, X, non-permission) with about 35% overhead on average and 57% in some cases. Besides, Occlum [168] proposed an MPX-based SFI technique for supporting in-enclave multi-processing, which adds a $10\times$ slowdown compared to Linux processes. SGXv2 adds special

instructions (e.g., EMODPR and EMODPE), to facilitate some of these issues including dynamic memory management. However, the overhead is not clear yet due to the lack of hardware support, and SGXv1 is still the dominating version for a while.

These improvements are not sufficient for most ML use cases. First, current designs can not be integrated with a wide range of accelerators such as GPUs, FPGAs, and TPUs. Despite recent efforts for supporting TEEs in these accelerators as described in Section IV-C, their systems software and ML framework designs and evaluations are in the early phases [169], [129], [142], [132]. Depending on the level of protection in different ML stages, the overhead and toolchain supports can be highly diverse. Also, none of the previous work offers scalable solutions with multiple accelerators.

C. Heterogeneity & Migration

Heterogeneous SoC architectures enable a wide range of functionalities and are now available even for modern IoT/edge platforms [170]. Modern SoCs contain heterogeneous CPUs (e.g., a combination of ARM and RISC-V architectures) and peripherals. As a result, the system stack on such devices includes multiple OSs (e.g., Linux and FreeRTOS), hypervisors, and different TEEs. As a simple example, recent NXP's i.MX boards, simultaneously support different TrustZone implementations for mixed cortex-A and cortex-M cores (e.g., i.MX 8ULP or 8QuadXPlus). Secure hardware partitioning and sharing hardware resources within such ever-growing complexity is challenging. That is the reason why most systems software (e.g., Android [171]) does not expose heterogeneous TEEs to userspace applications despite hardware availability.

Moreover, heterogeneous TEEs could have different security models and resource constraints which make the migration of TEE-based ML solutions from one platform to another much more challenging. For instance, remote attestation support in TrustZone-based and SGX-based solutions are different. Naively, combining different security protocols does not lead to a secure solution. Also, as summarized in Figure 4, vulnerabilities can easily be introduced and propagated in various

privilege boundaries (within the different or same address spaces). Detecting and protecting against such a complex attack surface is one of the primary obstacles to securely handling TEE migrations. Ad-hoc approaches are neither practical nor scale well. We need a principled approach to deal with heterogeneous TEEs and future security threats. A naive migration could widen or worsen the attack surface.

VII. NEXT GENERATION AND BEYOND

In this section, we present the potential approaches to tackle above discussed issues, towards better privacy-preserving machine learning atop TEEs.

A. Strengthen privacy foundation in Confidential Computing

As one TEE’s primary protection goals, privacy needs to be strictly defined and preferably be defined with theoretical basics (e.g., information theory). While most privacy measures are empirically evaluated by conducting attacks⁷, defining privacy theoretically is complicated due to the broad existence of attacks in ML.

Among candidates, *differential privacy* is a well-accepted definition of privacy. DP has been massively explored when applying to ML, but it faces the same generalization problem. As one example, one can apply DP-training [172] to show the protection efficiency against DRAs that recovers original training data from gradients computed on it. As expected, this protection is not reasonable because typical DP-training guarantees the “membership” privacy of the trained final model, not (and far from) intermediate gradients. To measure on the pre-gradient level, one needs to adjust the DP’s granularity from sample-level (i.e., sensitivity of samples in the dataset) to feature-level (i.e., sensitivity of features in the sample). This level of DP noises will highly likely destroy the model utility but potentially provides a theoretical-based sense of what is the level of privacy for such protection.

Another line of work is to apply computational constraint-based \mathcal{V} usable information theory [173]. \mathcal{V} -information allows us to consider an attack family with bound computational power and limited knowledge. With the defined attack family, one can measure the capability of the complete attack family, which may include many attacks, by computing \mathcal{V} -information similar to Shannon mutual information. This \mathcal{V} -information then reflects the privacy leakage [127]. Although such a method provides a theoretical base for measuring privacy, how to define a broad and general attack family still deserves further exploration.

Different types of privacy in ML are fragmental and overlapped, as concluded based on both previous empirical attacks and theoretical analyses. Without a solid understanding of privacy, TEE protection itself is not plausible; protection of one type of privacy could be invalid when the privacy goal changes. Therefore, a rigorous definition of privacy for specific use cases is the cornerstone for TEE protection.

⁷Measuring privacy using attack success rates fails to explain and quantify privacy leakages, and then leads to “Arms Race” where current protections needs re-evaluation or re-design when new attacks happen.

B. Partitioned ML execution for heterogeneous TEEs

For now and in the foreseeable future, most TEEs on modern devices will be heterogeneous, and many will keep small considering the ubiquitousness of devices and the requirement of small TCBs. It is desired to avoid monolithic ML frameworks; so we could partition the solution with a focus on protecting the most critical ML computation or stages.

The partition can be done on different scales. On a low level, most ML’s computation forms with numeral calculations e.g., Multiply-accumulate operations, floating point operations, or matrix multiplication. Running a portion of these computations inside TEE is straightforward, but determining whether computation is not an easy task. First, this determination needs well-defined privacy as above mentioned, e.g., which type of privacy to protect. Second, it requires the knowledge of how these computations have different leakage of the target privacy. Currently, we have a superficial understanding of layer-level privacy differences, i.e., the last layer contains the most membership privacy. However, there is still a need i) to deepen the layer-level understanding of more types of privacy; ii) to investigate efficient layer-wise partitioning approaches; iii) to investigate any more fine-grained privacy understanding (i.e., matrix multiplication and feature maps), as the layer-level partition is still heavy for constrained TEEs. Furthermore, one can also execute some kinds of assurance and verification inside TEEs instead of hiding sensitive parts directly. Such a method has been used for ML integrity checking [131] but still deserves further investigation i.e., combining with masking techniques to preserve privacy.

C. Dedicated TEE designs for general ML

Efficient ML execution requires modern computer architectures with parallelization features, e.g., multi-threading on CPUs to load data and on GPUs (or TPU, etc.) to perform matrix multiplication in forward and backward passes. These features are achieved by dedicatedly designing the hardware and the system of accelerated processing units, which most of current TEEs do not have. One solution is to equip TEEs with such parallelization features to increase the computational capability of ML. However, consequently, new features increase the TCB size. Synchronization bugs cause severe vulnerability of SGX [120], and one can also expect a considerable number of bugs that potentially exist in GPU-TEEs. How to avoid such a dilemma is really depending on how can we limit the trust boundary and reduce the TCB when applying parallel processing. One example is to consider one ML accelerated unit (GPU/TPU) to be physically isolated from the rest of the main-board and the processing system in the same way as TPM, and to be accessed only through a secure bus from TEE-enabled CPUs. Although this constrains the trust boundary with assumptions that this dedicated GPU serves only for such private ML and no one can physically breach the GPU, there are more use cases applicable such as MLaaS with proper remote attestation and verification.

Interestingly, another way is to avoid multi-core features but focus on in-memory computation. This potentially closes the

gap because the most recent TEE designs (SGXv2 and Arm CCA) support large memory sizes that will benefit in-memory ML in privacy-preserving manners [174], [175], [176].

D. TEE-aware ML

Due to the heterogeneity of devices' TEEs and ML workload, it is obvious that one solution cannot satisfy all needs. Therefore, designing the ML to be aware of TEEs and the capability of TEEs is critical. To achieve this, ML frameworks will determine the maximum workload to be executed with TEEs which first requires the TEE's information such as memory size, processor speed, protected storage, and potentially other features like multi-threading and secure communication channels. Then, ML computations from the most sensitive to less are deployed into TEEs. Such a determination involves many variables that influence factors from system cost to ML performance; thus, one can potentially form a multi-objective optimization problem similarly to Neural Architecture Search for resource-constrained environments [177], [178], [179]. Furthermore, one may need techniques such as sparse approach [180], [181] of knowledge distillation [182], [183] to handle the heterogeneity of TEEs and differences in model architectures, especially when updating models in FL.

E. Protecting the full ML pipeline

Protecting the full pipeline is intractable without the participation of multiple TEEs (e.g., [184], [22]). Multiple TEEs are able to provide multiple trusted areas covering more devices and consequently more pieces of the pipeline. To achieve such "multi-party computation" based on TEEs, one needs to enable a verification mechanism for multiple participants so that TEEs from different chip manufacturers can be remotely attested by two or more challengers. For example, potentially one third-party proxy attestation service can help to verify TEEs from both Arm and Intel [98], so that the TEE system could cover multiple locations of the ML pipeline.

In addition to the TEE workflow, determining the most sensitive parts of ML for protection is also not trivial. On the smallest scale, TEE provides a more trustworthy area in the ML pipeline, such as an additional trust base for random number producer or DP noise addition). Besides, investigations on the privacy or integrity of the pipeline's parts other than training/inference protection, such as data preparation, are vastly missing. Such protection on one specific part of the pipeline will involve again threat (privacy/integrity) definition, sophisticated TEE-based protection design, and performance measurement similar to the training stage protection. After such works on more stages of the pipeline, full ML pipeline protection will be formed concretely and be used to a greater extent and larger scale.

VIII. CONCLUSION

Protecting ML against privacy leakages and integrity breaches by the Confidential Computing techniques is an exciting and challenging new era. Although many studies have been dedicated to running the training and inference processes

inside TEEs, they still face the limitation of available trust resources. Since ML requires much more trusted resources; the current protection levels only provides the confidentiality and integrity of that specific stage, training/inference, in the complete ML pipeline. Current defenses against privacy leakage are usually measured through red-teaming and attack-specific metrics, which make the final results less reliable and generalized. Furthermore, large attack surfaces exist at stages, especially at the upstream stage like data preparation, which can cause huge negative impacts on the ML pipeline. Rethinking of how to tackle such a dilemma, i.e., the conflict between the large scale of the complete ML pipeline and the need for small TCB, is of utmost importance when conducting ML with Confidential Computing.

Confidential Computing achieves a hardware-based root-of-trust establishing a more trustworthy execution environment for ML activities, but we should rethink whether "hiding" the training/inference process inside such enclaves is the optimal solution. Future researchers/developers need to better understand the underlying privacy issues in the ML pipeline so that later protections focus on vital parts. The full ML pipeline needs to be evaluated to avoid fruitless labor on training/inference protection only. The current trends indicated that both TEE evolution and ML advances (i.e., partitioning ML execution for TEEs, and designing dedicated TEEs for ML) will lead to more *TEE-aware* future architectures. These potential solutions can help achieve much stronger privacy and security guarantees without compromising ML computational performance and introducing TEE's system overhead.

Suggestions to achieve ML with Confidential Computing for the hardware/system/framework designer are:

- Compared to security or integrity, privacy definition is still ambiguous. One needs to define the theoretical-based protection goal, e.g., with differential privacy or \mathcal{V} -usable information, to have well-grounded privacy guarantee.
- Protecting the upstream of the ML pipeline such as the data preparation is of utmost importance due to its lack has intractable and negative impacts. It can potentially be realized by integrating TEE-based verification into data signature. Multiple TEEs/Conclaves can further benefit the full ML pipeline protection.
- Developing dedicated TEEs for general ML requires rethinking and redesigning the hardware acceleration features and corresponding system stack. New paradigms such as in-memory computing that involves a smaller group of hardware and software modules may ease the condition due to the reduction of the TCB size and the attack surface.
- Designing the ML framework to be TEE-aware and partitionable for heterogeneous TEEs requires i) meticulously researching the privacy/integrity weakness of different ML granularity (from layer/feature map to numeral calculations), and ii) managing the TEE system to protect the most sensitive ML components efficiently (e.g., with masking) with a high priority.

REFERENCES

- [1] T. Orekondy, B. Schiele, and M. Fritz, “Knockoff nets: Stealing functionality of black-box models,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2019, pp. 4954–4963.
- [2] F. Tramèr, F. Zhang, A. Juels, M. K. Reiter, and T. Ristenpart, “Stealing machine learning models via prediction APIs,” in *25th {USENIX} Security Symposium ({USENIX} Security 16)*, 2016, pp. 601–618.
- [3] Y. Shen, X. He, Y. Han, and Y. Zhang, “Model stealing attacks against inductive graph neural networks,” in *2022 IEEE Symposium on Security and Privacy (SP)*. IEEE, 2022, pp. 1175–1192.
- [4] A. S. Rakin, M. H. I. Chowdhury, F. Yao, and D. Fan, “Deepsteal: Advanced model extractions leveraging efficient weight stealing in memories,” in *2022 IEEE Symposium on Security and Privacy (SP)*. IEEE, 2022, pp. 1157–1174.
- [5] M. Fredrikson, S. Jha, and T. Ristenpart, “Model inversion attacks that exploit confidence information and basic countermeasures,” in *Proceedings of the 22nd ACM SIGSAC conference on computer and communications security*, 2015, pp. 1322–1333.
- [6] Z. He, T. Zhang, and R. B. Lee, “Model inversion attacks against collaborative inference,” in *Proceedings of the 35th Annual Computer Security Applications Conference*, 2019, pp. 148–162.
- [7] M. Fang, X. Cao, J. Jia, and N. Gong, “Local model poisoning attacks to byzantine-robust federated learning,” in *29th {USENIX} Security Symposium ({USENIX} Security 20)*, 2020, pp. 1605–1622.
- [8] X. Chen, C. Liu, B. Li, K. Lu, and D. Song, “Targeted backdoor attacks on deep learning systems using data poisoning,” *arXiv preprint arXiv:1712.05526*, 2017.
- [9] V. Shejwalkar, A. Houmansadr, P. Kairouz, and D. Ramage, “Back to the drawing board: A critical evaluation of poisoning attacks on production federated learning,” in *2022 IEEE Symposium on Security and Privacy (SP)*. IEEE, 2022, pp. 1354–1371.
- [10] L. Zhu, Z. Liu, and S. Han, “Deep leakage from gradients,” *Advances in Neural Information Processing Systems*, vol. 32, pp. 14774–14784, 2019.
- [11] H. Yin, A. Mallya, A. Vahdat, J. M. Alvarez, J. Kautz, and P. Molchanov, “See through gradients: Image batch recovery via gradient inversion,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2021, pp. 16337–16346.
- [12] R. Shokri, M. Stronati, C. Song, and V. Shmatikov, “Membership inference attacks against machine learning models,” in *2017 IEEE Symposium on Security and Privacy (SP)*. IEEE, 2017, pp. 3–18.
- [13] L. Melis, C. Song, E. De Cristofaro, and V. Shmatikov, “Exploiting unintended feature leakage in collaborative learning,” in *2019 IEEE Symposium on Security and Privacy (SP)*. IEEE, 2019, pp. 691–706.
- [14] N. Carlini, S. Chien, M. Nasr, S. Song, A. Terzis, and F. Tramèr, “Membership inference attacks from first principles,” in *2022 IEEE Symposium on Security and Privacy (SP)*. IEEE, 2022, pp. 1897–1914.
- [15] I. J. Goodfellow, J. Shlens, and C. Szegedy, “Explaining and harnessing adversarial examples,” *arXiv preprint arXiv:1412.6572*, 2014.
- [16] C. C. Consortium, “A technical analysis of confidential computing v1.2.”
- [17] —, “Confidential computing: Hardware-based trusted execution for applications and data.”
- [18] O. Ohrimenko, F. Schuster, C. Fournet, A. Mehta, S. Nowozin, K. Vaswani, and M. Costa, “Oblivious multi-party machine learning on trusted processors,” in *25th {USENIX} Security Symposium ({USENIX} Security 16)*, 2016, pp. 619–636.
- [19] F. Tramèr and D. Boneh, “Slalom: Fast, verifiable and private execution of neural networks in trusted hardware,” in *International Conference on Learning Representations*, 2018.
- [20] T. Hunt, C. Song, R. Shokri, V. Shmatikov, and E. Witchel, “Chiron: Privacy-preserving machine learning as a service,” *arXiv preprint arXiv:1803.05961*, 2018.
- [21] T. Lee, Z. Lin, S. Pushp, C. Li, Y. Liu, Y. Lee, F. Xu, C. Xu, L. Zhang, and J. Song, “Occlumency: Privacy-preserving remote deep-learning inference using SGX,” in *The 25th Annual International Conference on Mobile Computing and Networking*, 2019, pp. 1–17.
- [22] F. Mo, H. Haddadi, K. Katevas, E. Marin, D. Perino, and N. Kourtellis, “PPFL: privacy-preserving federated learning with trusted execution environments,” *arXiv preprint arXiv:2104.14380*, 2021.
- [23] “Azure confidential computing,” <https://azure.microsoft.com/en-gb/solutions/confidential-compute/#overview>, accessed: 2022-02-03.
- [24] “AWS Nitro Enclaves,” <https://aws.amazon.com/ec2/nitro/nitro-enclaves/>, accessed: 2022-02-03.
- [25] “Google Cloud Confidential Computing,” <https://cloud.google.com/confidential-computing>, accessed: 2022-02-03.
- [26] K. D. Duy, T. Noh, S. Huh, and H. Lee, “Confidential machine learning computation in untrusted environments: A systems security perspective,” *IEEE Access*, vol. 9, pp. 168 656–168 677, 2021.
- [27] M. I. Jordan and T. M. Mitchell, “Machine learning: Trends, perspectives, and prospects,” *Science*, vol. 349, no. 6245, pp. 255–260, 2015.
- [28] Y. LeCun, Y. Bengio, and G. Hinton, “Deep learning,” *nature*, vol. 521, no. 7553, pp. 436–444, 2015.
- [29] X. Zhu and A. B. Goldberg, “Introduction to semi-supervised learning,” *Synthesis lectures on artificial intelligence and machine learning*, vol. 3, no. 1, pp. 1–130, 2009.
- [30] D. Berthelot, N. Carlini, I. Goodfellow, N. Papernot, A. Oliver, and C. A. Raffel, “Mixmatch: A holistic approach to semi-supervised learning,” *Advances in Neural Information Processing Systems*, vol. 32, 2019.
- [31] Q. V. Le, “Building high-level features using large scale unsupervised learning,” in *2013 IEEE international conference on acoustics, speech and signal processing*. IEEE, 2013, pp. 8595–8598.
- [32] R. S. Sutton and A. G. Barto, *Reinforcement learning: An introduction*. MIT press, 2018.
- [33] C. Shorten and T. M. Khoshgoufar, “A survey on image data augmentation for deep learning,” *Journal of Big Data*, vol. 6, no. 1, pp. 1–48, 2019.
- [34] Q. Xie, Z. Dai, E. Hovy, T. Luong, and Q. Le, “Unsupervised data augmentation for consistency training,” *Advances in Neural Information Processing Systems*, vol. 33, 2020.
- [35] Y. LeCun, L. Bottou, Y. Bengio, and P. Haffner, “Gradient-based learning applied to document recognition,” *Proceedings of the IEEE*, vol. 86, no. 11, pp. 2278–2324, 1998.
- [36] L. Bottou, “Large-scale machine learning with stochastic gradient descent,” in *Proceedings of COMPSTAT’2010*. Springer, 2010, pp. 177–186.
- [37] R. Kohavi *et al.*, “A study of cross-validation and bootstrap for accuracy estimation and model selection,” in *International Joint Conference on Artificial Intelligence*, vol. 14, no. 2. Montreal, Canada, 1995, pp. 1137–1145.
- [38] S. Arlot and A. Celisse, “A survey of cross-validation procedures for model selection,” *Statistics surveys*, vol. 4, pp. 40–79, 2010.
- [39] S. J. Pan and Q. Yang, “A survey on transfer learning,” *IEEE Transactions on knowledge and data engineering*, vol. 22, no. 10, pp. 1345–1359, 2009.
- [40] P. Marcelino, “Transfer learning from pre-trained models,” *Towards Data Science*, 2018.
- [41] J. Yosinski, J. Clune, Y. Bengio, and H. Lipson, “How transferable are features in deep neural networks?” *arXiv preprint arXiv:1411.1792*, 2014.
- [42] M. D. Zeiler and R. Fergus, “Visualizing and understanding convolutional networks,” in *European conference on computer vision*. Springer, 2014, pp. 818–833.
- [43] H. Qi, M. Brown, and D. G. Lowe, “Low-shot learning with imprinted weights,” in *Proceedings of the IEEE conference on computer vision and pattern recognition*, 2018, pp. 5822–5830.
- [44] B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. A. y Arcas, “Communication-efficient learning of deep networks from decentralized data,” in *Artificial intelligence and statistics*. PMLR, 2017, pp. 1273–1282.
- [45] T. Li, A. K. Sahu, A. Talwalkar, and V. Smith, “Federated learning: Challenges, methods, and future directions,” *IEEE Signal Processing Magazine*, vol. 37, no. 3, pp. 50–60, 2020.
- [46] P. Kairouz, H. B. McMahan, B. Avent, A. Bellet, M. Bennis, A. N. Bhagoji, K. Bonawitz, Z. Charles, G. Cormode, R. Cummings *et al.*, “Advances and open problems in federated learning,” *arXiv preprint arXiv:1912.04977*, 2019.
- [47] M. Abadi, P. Barham, J. Chen, Z. Chen, A. Davis, J. Dean, M. Devin, S. Ghemawat, G. Irving, M. Isard *et al.*, “Tensorflow: A system for large-scale machine learning,” in *12th {USENIX} symposium on operating systems design and implementation ({OSDI} 16)*, 2016, pp. 265–283.

- [48] Y. Yao, Z. Xiao, B. Wang, B. Viswanath, H. Zheng, and B. Y. Zhao, "Complexity vs. performance: empirical analysis of machine learning as a service," in *Proceedings of the 2017 Internet Measurement Conference*, 2017, pp. 384–397.
- [49] M. Li, D. G. Andersen, J. W. Park, A. J. Smola, A. Ahmed, V. Josifovski, J. Long, E. J. Shekita, and B.-Y. Su, "Scaling distributed machine learning with the parameter server," in *11th {USENIX} Symposium on Operating Systems Design and Implementation ({OSDI} 14)*, 2014, pp. 583–598.
- [50] J. Geiping, H. Bauermeister, H. Dröge, and M. Moeller, "Inverting gradients—how easy is it to break privacy in federated learning?" *arXiv preprint arXiv:2003.14053*, 2020.
- [51] M. Ribeiro, K. Grolinger, and M. A. Capretz, "Mlaas: Machine learning as a service," in *2015 IEEE 14th International Conference on Machine Learning and Applications (ICMLA)*. IEEE, 2015, pp. 896–902.
- [52] E. Hesamifard, H. Takabi, M. Ghasemi, and R. N. Wright, "Privacy-preserving machine learning as a service," *Proceedings on Privacy Enhancing Technologies*, vol. 3, pp. 123–142, 2018.
- [53] M. Nasr, R. Shokri, and A. Houmansadr, "Comprehensive privacy analysis of deep learning: Passive and active white-box inference attacks against centralized and federated learning," in *2019 IEEE symposium on security and privacy (SP)*. IEEE, 2019, pp. 739–753.
- [54] B. Hitaj, G. Ateniese, and F. Perez-Cruz, "Deep models under the gan: information leakage from collaborative deep learning," in *Proceedings of the 2017 ACM SIGSAC Conference on Computer and Communications Security*, 2017, pp. 603–618.
- [55] J. Jia and N. Z. Gong, "Attriguard: A practical defense against attribute inference attacks via adversarial machine learning," in *27th {USENIX} Security Symposium ({USENIX} Security 18)*, 2018, pp. 513–529.
- [56] H. Yu, K. Yang, T. Zhang, Y.-Y. Tsai, T.-Y. Ho, and Y. Jin, "Cloudeleak: Large-scale deep learning models stealing through adversarial examples," in *Network and Distributed System Security Symposium (NDSS)*, 2020.
- [57] S.-M. Moosavi-Dezfooli, A. Fawzi, O. Fawzi, and P. Frossard, "Universal adversarial perturbations," in *Proceedings of the IEEE conference on computer vision and pattern recognition*, 2017, pp. 1765–1773.
- [58] Z. Tarkhani, L. Qendro, M. O. Brown, O. Hill, C. Mascolo, and A. Madhavapeddy, "Enhancing the security & privacy of wearable brain-computer interfaces," *arXiv preprint arXiv:2201.07711*, 2022.
- [59] J. R. Sanchez Vicarte, B. Schreiber, R. Paccagnella, and C. W. Fletcher, "Game of threads: Enabling asynchronous poisoning attacks," in *Proceedings of the Twenty-Fifth International Conference on Architectural Support for Programming Languages and Operating Systems*, 2020, pp. 35–52.
- [60] V. Tolpegin, S. Truex, M. E. Gursoy, and L. Liu, "Data poisoning attacks against federated learning systems," in *European Symposium on Research in Computer Security*. Springer, 2020, pp. 480–501.
- [61] Y. Liu, X. Ma, J. Bailey, and F. Lu, "Reflection backdoor: A natural backdoor attack on deep neural networks," in *European Conference on Computer Vision*. Springer, 2020, pp. 182–199.
- [62] E. Bagdasaryan, A. Veit, Y. Hua, D. Estrin, and V. Shmatikov, "How to backdoor federated learning," in *International Conference on Artificial Intelligence and Statistics*. PMLR, 2020, pp. 2938–2948.
- [63] J. Jia, Y. Liu, and N. Z. Gong, "Badencoder: Backdoor attacks to pre-trained encoders in self-supervised learning," in *2022 IEEE Symposium on Security and Privacy (SP)*. IEEE, 2022, pp. 2043–2059.
- [64] N. Papernot, P. McDaniel, A. Sinha, and M. P. Wellman, "SoK: Security and privacy in machine learning," in *2018 IEEE European Symposium on Security and Privacy (EuroS&P)*. IEEE, 2018, pp. 399–414.
- [65] A. Sablayrolles, M. Douze, C. Schmid, Y. Ollivier, and H. Jégou, "White-box vs black-box: Bayes optimal strategies for membership inference," in *International Conference on Machine Learning*. PMLR, 2019, pp. 5558–5567.
- [66] S. Yeom, I. Giacomelli, M. Fredrikson, and S. Jha, "Privacy risk in machine learning: Analyzing the connection to overfitting," in *2018 IEEE 31st computer security foundations symposium (CSF)*. IEEE, 2018, pp. 268–282.
- [67] T. L. F. Project. Confidential computing consortium. [Online]. Available: <https://confidentialcomputing.io/>
- [68] GlobalPlatform, *White Paper on the Trusted Execution Environment*. GlobalPlatform, 2015. [Online]. Available: https://globalplatform.org/wp-content/uploads/2018/04/GlobalPlatform_TEE_Whitepaper_2015.pdf
- [69] L. Guan, P. Liu, X. Xing, X. Ge, S. Zhang, M. Yu, and T. Jaeger, "Trustshadow: Secure execution of unmodified applications with arm trustzone," in *Proceedings of the 15th Annual International Conference on Mobile Systems, Applications, and Services*. ACM, 2017, pp. 488–501.
- [70] R. Bahmani, F. Brasser, G. Dessouky, P. Jauernig, M. Klimmek, A.-R. Sadeghi, and E. Stappf, "{CURE}: A security architecture with {Customizable} and resilient enclaves," in *30th USENIX Security Symposium (USENIX Security 21)*, 2021, pp. 1073–1090.
- [71] J. Lee, J. Jang, Y. Jang, N. Kwak, Y. Choi, C. Choi, T. Kim, M. Peinado, and B. B. Kang, "Hacking in Darkness: Return-Oriented Programming against Secure Enclaves," in *Proceedings of the 26th USENIX Conference on Security Symposium*, 2017, pp. 523–539.
- [72] M. R. Khandaker, Y. Cheng, Z. Wang, and T. Wei, "Coin attacks: On insecurity of enclave untrusted interfaces in sgx," in *Proceedings of the Twenty-Fifth International Conference on Architectural Support for Programming Languages and Operating Systems*, 2020, pp. 971–985.
- [73] Van Bulck, Jo and Oswald, David and Marin, Eduard and Aldoseri, Abdulla and Garcia, Flavio D. and Piessens, Frank, "A Tale of Two Worlds: Assessing the Vulnerability of Enclave Shielding Runtimes," in *Proceedings of the ACM SIGSAC Conference on Computer and Communications Security (CCS)*, 2019, pp. 1741–1758.
- [74] Z. Chen, G. Vasilakis, K. Murdock, E. Dean, D. Oswald, and F. D. Garcia, "Voltpillager: Hardware-based fault injection attacks against intel SGX enclaves using the SVID voltage scaling interface," in *30th USENIX Security Symposium*, Vancouver, B.C., Aug. 2021.
- [75] M. Li, L. Wilke, J. Wichelmann, T. Eisenbarth, R. Teodorescu, and Y. Zhang, "A systematic look at ciphertext side channels on amd sev-snp," in *2022 IEEE Symposium on Security and Privacy (SP)*. IEEE Computer Society, 2022, pp. 1541–1541.
- [76] M. Lipp, A. Kogler, D. Oswald, M. Schwarz, C. Easdon, C. Canella, and D. Gruss, "PLATYPUS: Software-based Power Side-Channel Attacks on x86," in *2021 IEEE Symposium on Security and Privacy (SP)*. IEEE, 2021.
- [77] A. Machiry, E. Gustafson, C. Spensky, C. Salls, N. Stephens, R. Wang, A. Bianchi, Y. R. Choe, C. Kruegel, and G. Vigna, "Boomerang: Exploiting the semantic gap in trusted execution environments," in *NDSS*, 2017.
- [78] D. Cerdeira, N. Santos, P. Fonseca, and S. Pinto, "SoK: Understanding the prevailing security vulnerabilities in trustzone-assisted tee systems," in *2020 IEEE Symposium on Security and Privacy (SP)*. IEEE, 2020, pp. 1416–1432.
- [79] H. Xia, D. Zhang, W. Liu, I. Haller, B. Sherwin, and D. Chisnall, "A secret-free hypervisor: Rethinking isolation in the age of speculative vulnerabilities," in *2022 IEEE Symposium on Security and Privacy (SP)*. IEEE Computer Society, 2022, pp. 1544–1544.
- [80] V. Costan and S. Devadas, "Intel SGX Explained," *IACR Cryptol. ePrint Arch.*, vol. 2016, no. 86, pp. 1–118, 2016.
- [81] V. Haldar, D. Chandra, and M. Franz, "Semantic remote attestation: A virtual machine directed approach to trusted computing," in *USENIX Virtual Machine Research and Technology Symposium*, vol. 2004, 2004.
- [82] G. Coker, J. Guttman, P. Loscocco, A. Herzog, J. Millen, B. O'Hanlon, J. Ramsdell, A. Segall, J. Sheehy, and B. Sniffen, "Principles of remote attestation," *International Journal of Information Security*, vol. 10, no. 2, pp. 63–81, 2011.
- [83] V. Scarlata, S. Johnson, J. Beaney, and P. Zmijewski, "Supporting third party attestation for intel sgx with intel data center attestation primitives," *White paper*, 2018.
- [84] S. Johnson, V. Scarlata, C. Rozas, E. Brickell, and F. Mckeen, "Intel software guard extensions: Epid provisioning and attestation services," *White Paper*, vol. 1, no. 1-10, p. 119, 2016.
- [85] "Intel® trust domain extensions (intel® tdx)," <https://software.intel.com/content/www/us/en/develop/articles/intel-trust-domain-extensions.html>, 2020.
- [86] Microsoft, "Ms-iot ftpm," <https://github.com/microsoft/ms-tpm-20-ref/blob/master/Samples/ARM32-Firmware/TPM/README.md>.
- [87] "Intel Software Guard Extensions SDK," <https://01.org/intel-software-guard-extensions>, accessed: 2021-10-21.
- [88] "Open Portable Trusted Execution Environment," <https://www.op-tee.org/>, accessed: 2021-10-21.
- [89] "MultiZone Security for RISC-V," <https://hex-five.com/multizone-security-sdk/>, accessed: 2021-10-21.
- [90] "Keystone: An Open Framework for Architecting Trusted Execution Environments," <https://keystone-enclave.org/>, accessed: 2021-10-21.

- [91] "Open Enclave SDK," <https://openenclave.io/sdk/>, accessed: 2021-10-21.
- [92] "Asylo: An open and flexible framework for enclave applications," <https://asylo.dev/>, accessed: 2021-10-21.
- [93] "Apache Teaclave (incubating)," <https://teaclave.apache.org/>, accessed: 2021-10-21.
- [94] "Gramine - a Library OS for Unmodified Applications," <https://grapheneproject.io/>, accessed: 2021-10-21.
- [95] "Occlum - A library OS empowering everyone to run every application in secure enclaves," <https://occlum.io/>, accessed: 2021-10-21.
- [96] "SCONE - A secure Container Environment," <https://scontain.com/>, accessed: 2021-10-21.
- [97] "Enarx," <https://enarx.dev/>, accessed: 2021-10-21.
- [98] "Veracruz: Privacy-Preserving Collaborative Compute," <https://veracruz-project.com/>, accessed: 2021-10-21.
- [99] A. Baumann, M. Peinado, and G. Hunt, "Shielding applications from an untrusted cloud with Haven," *ACM Transactions on Computer Systems (TOCS)*, vol. 33, no. 3, p. 8, 2015.
- [100] D. E. Porter, S. Boyd-Wickizer, J. Howell, R. Olinsky, and G. C. Hunt, "Rethinking the library os from the top down," in *ACM SIGPLAN Notices*, vol. 46. ACM, 2011, pp. 291–304.
- [101] C.-C. Tsai, D. E. Porter, and M. Vij, "Graphene-sgx: A practical library os for unmodified applications on sgx," in *Proceedings of the USENIX Annual Technical Conference (ATC)*, 2017, p. 8.
- [102] S. Arnaudov, B. Trach, F. Gregor, T. Knauth, A. Martin, C. Priebe, J. Lind, D. Muthukumaran, D. O'keeffe, M. Stillwell *et al.*, "Scone: Secure linux containers with intel sgx," in *OSDI*, vol. 16, 2016, pp. 689–703.
- [103] I. C. London. (2019) Sgx-ikl library os for running linux applications inside of intel sgx enclaves. <https://github.com/llds/sgx-ikl>. Access Date :2019-10-01.
- [104] I. Corporation. (2019) Intel(r) software guard extensions for linux os. <https://github.com/intel/linux-sgx>. Access Date :2019-03-01.
- [105] M. Corporation, "Open enclave sdk," <https://github.com/openenclave/openenclave>, 2019, access Date :2019-08-12.
- [106] Google, "Asylo - an open and flexible framework for enclave applications," <http://web.archive.org/web/20080207010024/http://www.808multimedia.com/winnt/kernel.htm>, 2018.
- [107] D. Lee, D. Kohlbrenner, S. Shinde, K. Asanovic, and D. Song, "Keystone: An open framework for architecting trusted execution environments," in *Proceedings of the Fifteenth European Conference on Computer Systems*, ser. EuroSys '20, 2020.
- [108] C.-C. Tsai, J. Son, B. Jain, J. McAvey, R. A. Popa, and D. E. Porter, "Civet: An efficient java partitioning framework for hardware enclaves," in *29th {USENIX} Security Symposium ({USENIX} Security 20)*, 2020.
- [109] N. Santos, H. Raj, S. Saroui, and A. Wolman, "Using arm trustzone to build a trusted language runtime for mobile applications," *ACM SIGARCH Computer Architecture News*, vol. 42, no. 1, pp. 67–80, 2014.
- [110] J. Lind, C. Priebe, D. Muthukumaran, D. O'Keeffe, P.-L. Aublin, F. Kelbert, T. Reiher, D. Goltzsche, D. Eysers, R. Kapitza *et al.*, "Glamdring: Automatic application partitioning for intel {SGX}," in *2017 USENIX Annual Technical Conference (USENIX ATC 17)*, 2017, pp. 285–298.
- [111] Qualcomm, *Guard Your Data with the Qualcomm Snapdragon Mobile Platform*. Qualcomm, 2019. [Online]. Available: <https://www.qualcomm.com/media/documents/files/guard-your-data-with-the-qualcomm-snapdragon-mobile-platform.pdf>
- [112] Trustonic. Trustonic: World leading embedded cybersecurity technology trustonic. [Online]. Available: <https://www.trustonic.com/>
- [113] M. Busch, J. Westphal, and T. Mueller, "Unearthing the trustedcore: A critical review on huawei's trusted execution environment," in *14th {USENIX} Workshop on Offensive Technologies ({WOOT} 20)*, 2020.
- [114] K. Grover, S. Tople, S. Shinde, R. Bhagwan, and R. Ramjee, "Privado: Practical and secure dnn inference with enclaves," *arXiv preprint arXiv:1810.00602*, 2018.
- [115] N. Hynes, R. Cheng, and D. Song, "Efficient deep learning on multi-source private data," *arXiv preprint arXiv:1807.06689*, 2018.
- [116] S. P. Bayerl, T. Frassetto, P. Jauernig, K. Riedhammer, A.-R. Sadeghi, T. Schneider, E. Stapf, and C. Weinert, "Offline Model Guard: Secure and private ML on mobile devices," in *2020 Design, Automation & Test in Europe Conference & Exhibition (DATE)*. IEEE, 2020, pp. 460–465.
- [117] F. Mo, A. S. Shamsabadi, K. Katevas, S. Demetriou, I. Leontiadis, A. Cavallaro, and H. Haddadi, "DarkneTZ: towards model privacy at the edge using trusted execution environments," in *Proceedings of the 18th International Conference on Mobile Systems, Applications, and Services*, 2020, pp. 161–174.
- [118] F. McKeen, I. Alexandrovich, I. Anati, D. Caspi, S. Johnson, R. Leslie-Hurd, and C. Rozas, "Intel® software guard extensions (intel® sgx) support for dynamic memory management inside an enclave," in *Proceedings of the Hardware and Architectural Support for Security and Privacy 2016*, 2016, pp. 1–9.
- [119] R. Kunkel, D. L. Quoc, F. Gregor, S. Arnaudov, P. Bhatotia, and C. Fetzer, "TensorSCONE: A secure tensorflow framework using Intel SGX," *arXiv preprint arXiv:1902.04413*, 2019.
- [120] N. Weichbrodt, A. Kurmus, P. Pietzuch, and R. Kapitza, "Asyncshock: Exploiting synchronisation bugs in intel sgx enclaves," in *European Symposium on Research in Computer Security*. Springer, 2016, pp. 440–457.
- [121] R. Asakura. rust-autograd. [Online]. Available: <https://github.com/rask/rust-autograd>
- [122] A. Paszke, S. Gross, F. Massa, A. Lerer, J. Bradbury, G. Chanan, T. Killeen, Z. Lin, N. Gimelshein, L. Antiga *et al.*, "PyTorch: An imperative style, high-performance deep learning library," *Advances in neural information processing systems*, vol. 32, pp. 8026–8037, 2019.
- [123] J. Jia, A. Salem, M. Backes, Y. Zhang, and N. Z. Gong, "Memguard: Defending against black-box membership inference attacks via adversarial examples," in *Proceedings of the 2019 ACM SIGSAC conference on computer and communications security*, 2019, pp. 259–274.
- [124] S. E. Oh, T. Yang, N. Mathews, J. K. Holland, M. S. Rahman, N. Hopper, and M. Wright, "Deepcoffea: Improved flow correlation attacks on tor via metric learning and amplification," in *2022 IEEE Symposium on Security and Privacy (SP)*. IEEE, 2022, pp. 1915–1932.
- [125] Z. Yang, Z. Yuan, S. Jin, X. Chen, L. Sun, X. Du, W. Li, and H. Zhang, "Fsaflow: Lightweight and fast dynamic path tracking and control for privacy protection on android using hybrid analysis with state-reduction strategy," in *Proceedings of the 43rd IEEE Symposium on Security and Privacy (SP)*, IEEE, San Francisco, CA, USA, 2022, pp. 23–25.
- [126] F. Mo, A. Borovykh, M. Malekzadeh, H. Haddadi, and S. Demetriou, "Layer-wise characterization of latent information leakage in federated learning," *arXiv preprint arXiv:2010.08762*, 2020.
- [127] —, "Quantifying information leakage from gradients," *arXiv preprint arXiv:2105.13929*, 2021.
- [128] Z. Gu, H. Huang, J. Zhang, D. Su, A. Lamba, D. Pendarakis, and I. Molloy, "Securing input data of deep learning inference systems via partitioned enclave execution," *arXiv preprint arXiv:1807.00969*, 2018.
- [129] S. Volos, K. Vaswani, and R. Bruno, "Graviton: Trusted execution environments on GPUs," in *13th {USENIX} Symposium on Operating Systems Design and Implementation ({OSDI} 18)*, 2018, pp. 681–696.
- [130] Z. Gu, H. Huang, J. Zhang, D. Su, H. Jamjoom, A. Lamba, D. Pendarakis, and I. Molloy, "Confidential inference via ternary model partitioning," *arXiv preprint arXiv:1807.00969*, 2018.
- [131] X. Zhang, F. Li, Z. Zhang, Q. Li, C. Wang, and J. Wu, "Enabling execution assurance of federated learning at untrusted participants," in *IEEE INFOCOM 2020-IEEE Conference on Computer Communications*. IEEE, 2020, pp. 1877–1886.
- [132] T. Hunt, Z. Jia, V. Miller, A. Szekely, Y. Hu, C. J. Rossbach, and E. Witchel, "Telekine: Secure computing with cloud gpus," in *17th {USENIX} Symposium on Networked Systems Design and Implementation ({NSDI} 20)*, 2020, pp. 817–833.
- [133] L. Hanzlik, Y. Zhang, K. Grosse, A. Salem, M. Augustin, M. Backes, and M. Fritz, "MLcapsule: Guarded offline deployment of machine learning as a service," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2021, pp. 3300–3309.
- [134] Z. Liu, Y. Lu, X. Xie, Y. Fang, Z. Jian, and T. Li, "Trusted-DNN: A trustzone-based adaptive isolation strategy for deep neural networks," in *ACM Turing Award Celebration Conference-China (ACM TURC 2021)*, 2021, pp. 67–71.
- [135] J.-B. Truong, W. Gallagher, T. Guo, and R. J. Walls, "Memory-efficient deep learning inference in trusted execution environments," *arXiv preprint arXiv:2104.15109*, 2021.
- [136] L. K. Ng, S. S. Chow, A. P. Woo, D. P. Wong, and Y. Zhao, "Goten: Gpu-outsourcing trusted execution of neural network training," in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 35, no. 17, 2021, pp. 14 876–14 883.

- [137] C. Zhang, J. Xia, B. Yang, H. Puyang, W. Wang, R. Chen, I. E. Akkus, P. Aditya, and F. Yan, "Citadel: Protecting data privacy and model confidentiality for collaborative learning," in *Proceedings of the ACM Symposium on Cloud Computing*, 2021, pp. 546–561.
- [138] S. Park, S. Kim, and Y.-s. Lim, "Fairness audit of machine learning models with confidential computing," in *Proceedings of the ACM Web Conference 2022*, 2022, pp. 3488–3499.
- [139] F. Mo, H. Haddadi, K. Katevas, E. Marin, D. Perino, and N. Kourtellis, "Ppfl: Enhancing privacy in federated learning with confidential computing," *GetMobile: Mobile Computing and Communications*, vol. 25, no. 4, pp. 35–38, 2022.
- [140] A. Krizhevsky, I. Sutskever, and G. E. Hinton, "Imagenet classification with deep convolutional neural networks," *Advances in neural information processing systems*, vol. 25, pp. 1097–1105, 2012.
- [141] M. Zhao, M. Gao, and C. Kozryakis, "Shef: shielded enclaves for cloud fpgas," in *Proceedings of the 27th ACM International Conference on Architectural Support for Programming Languages and Operating Systems*, 2022, pp. 1070–1085.
- [142] K. Xia, Y. Luo, X. Xu, and S. Wei, "Sgx-fpga: Trusted execution environment for cpu-fpga heterogeneous architecture," in *2021 58th ACM/IEEE Design Automation Conference (DAC)*. IEEE, 2021, pp. 301–306.
- [143] H. Park and F. X. Lin, "Safe and practical gpu acceleration in trustzone," *arXiv preprint arXiv:2111.03065*, 2021.
- [144] NVIDIA, "Nvidia h100 tensor core gpu architecture."
- [145] X. Li, H. Xiong, X. Li, X. Wu, X. Zhang, J. Liu, J. Bian, and D. Dou, "Interpretable deep learning: Interpretation, interpretability, trustworthiness, and beyond," *arXiv preprint arXiv:2103.10689*, 2021.
- [146] F. Tramèr, A. Kurakin, N. Papernot, I. Goodfellow, D. Boneh, and P. McDaniel, "Ensemble adversarial training: Attacks and defenses," in *International Conference on Learning Representations*, 2018.
- [147] A. R. Chowdhury, C. Guo, S. Jha, and L. van der Maaten, "Eiffel: Ensuring integrity for federated learning," *arXiv preprint arXiv:2112.12727*, 2021.
- [148] R. Motwani and P. Raghavan, "Randomized algorithms," *ACM Computing Surveys (CSUR)*, vol. 28, no. 1, pp. 33–37, 1996.
- [149] R. Zhang, J. Liu, Y. Ding, Z. Wang, Q. Wu, and K. Ren, "adversarial examples" for proof-of-learning," in *2022 IEEE Symposium on Security and Privacy (SP)*. IEEE, 2022, pp. 1408–1422.
- [150] R. C. Merkle, "Protocols for public key cryptosystems," in *1980 IEEE Symposium on Security and Privacy*. IEEE, 1980, pp. 122–122.
- [151] N. Lukas, E. Jiang, X. Li, and F. Kerschbaum, "Sok: How robust is image classification deep neural network watermarking?" in *2022 IEEE Symposium on Security and Privacy (SP)*. IEEE, 2022, pp. 787–804.
- [152] R. C. Merkle, "A certified digital signature," in *Conference on the Theory and Application of Cryptology*. Springer, 1989, pp. 218–238.
- [153] F. Cohen, "A cryptographic checksum for integrity protection," *Computers & Security*, vol. 6, no. 6, pp. 505–510, 1987.
- [154] N. Karapanos, A. Filios, R. A. Popa, and S. Capkun, "Verena: End-to-end integrity protection for web applications," in *2016 IEEE Symposium on Security and Privacy (SP)*. IEEE, 2016, pp. 895–913.
- [155] C. E. Brodley and M. A. Friedl, "Identifying mislabeled training data," *Journal of artificial intelligence research*, vol. 11, pp. 131–167, 1999.
- [156] X. Chu, I. F. Ilyas, S. Krishnan, and J. Wang, "Data cleaning: Overview and emerging challenges," in *Proceedings of the 2016 international conference on management of data*, 2016, pp. 2201–2206.
- [157] D. Bahri, H. Jiang, and M. Gupta, "Deep k-nn for noisy labels," in *International Conference on Machine Learning*. PMLR, 2020, pp. 540–550.
- [158] Z. Tarkhani and A. Madhavapeddy, "Enclave-aware compartmentalization and secure sharing with sirius," *arXiv preprint arXiv:2009.01869*, 2020.
- [159] D. Suci, S. McLaughlin, L. Simon, and R. Sion, "Horizontal privilege escalation in trusted applications," in *29th {USENIX} Security Symposium ({USENIX} Security 20)*, 2020.
- [160] S. Checkoway and H. Shacham, *Iago attacks: why the system call API is a bad untrusted RPC interface*. ACM, 2013, vol. 41, no. 1.
- [161] J. Van Bulck, D. Oswald, E. Marin, A. Aldoseri, F. D. Garcia, and F. Piessens, "A tale of two worlds: Assessing the vulnerability of enclave shielding runtimes," in *Proceedings of the 2019 ACM SIGSAC Conference on Computer and Communications Security*, 2019, pp. 1741–1758.
- [162] S. Weiser, L. Mayr, M. Schwarz, and D. Gruss, "Sgxjail: Defeating enclave malware via confinement," in *22nd International Symposium on Research in Attacks, Intrusions and Defenses (RAID 2019)*, 2019, pp. 353–366.
- [163] Y. Liu and A. Srivastava, "Ganred: Gan-based reverse engineering of dnns via cache side-channel," in *Proceedings of the 2020 ACM SIGSAC Conference on Cloud Computing Security Workshop*, 2020, pp. 41–52.
- [164] Y. Xiang, Z. Chen, Z. Chen, Z. Fang, H. Hao, J. Chen, Y. Liu, Z. Wu, Q. Xuan, and X. Yang, "Open dnn box by power side-channel attack," *IEEE Transactions on Circuits and Systems II: Express Briefs*, vol. 67, no. 11, pp. 2717–2721, 2020.
- [165] L. Batina, S. Bhasin, D. Jap, and S. Picek, "{CSI}-{NN}: Reverse engineering of neural network architectures through electromagnetic side channel," in *28th USENIX Security Symposium (USENIX Security 19)*, 2019, pp. 515–532.
- [166] W. Zheng, A. Dave, J. G. Beekman, R. A. Popa, J. E. Gonzalez, and I. Stoica, "Opaque: An oblivious and encrypted distributed analytics platform," in *14th USENIX Symposium on Networked Systems Design and Implementation (NSDI 17)*, 2017, pp. 283–298.
- [167] W. Zhao, K. Lu, Y. Qi, and S. Qi, "Mptec: bringing flexible and efficient memory protection to intel sgx," in *Proceedings of the Fifteenth European Conference on Computer Systems*, 2020, pp. 1–15.
- [168] Y. Shen, H. Tian, Y. Chen, K. Chen, R. Wang, Y. Xu, Y. Xia, and S. Yan, "Occlum: Secure and efficient multitasking inside a single enclave of intel sgx," in *Proceedings of the Twenty-Fifth International Conference on Architectural Support for Programming Languages and Operating Systems*, 2020, pp. 955–970.
- [169] I. Jang, A. Tang, T. Kim, S. Sethumadhavan, and J. Huh, "Heterogeneous isolated execution for commodity gpus," in *Proceedings of the Twenty-Fourth International Conference on Architectural Support for Programming Languages and Operating Systems*, 2019, pp. 455–468.
- [170] S. Zhao, M. Li, Y. Zhang, and Z. Lin, "vsgx: Virtualizing sgx enclaves on amd sev," in *2022 IEEE Symposium on Security and Privacy (SP)*. IEEE, 2022, pp. 321–336.
- [171] NXP. Asug-i.mx android security user's guide. [Online]. Available: https://www.nxp.com/docs/en/user-guide/IMX_ANDROID_SECURITY_USERS_GUIDE.pdf
- [172] M. Abadi, A. Chu, I. Goodfellow, H. B. McMahan, I. Mironov, K. Talwar, and L. Zhang, "Deep learning with differential privacy," in *Proceedings of the 2016 ACM SIGSAC conference on computer and communications security*, 2016, pp. 308–318.
- [173] Y. Xu, S. Zhao, J. Song, R. Stewart, and S. Ermon, "A theory of usable information under computational constraints," in *International Conference on Learning Representations*, 2019.
- [174] E. Eleftheriou, M. Le Gallo, S. Nandakumar, C. Piveteau, I. Boybat, V. Joshi, R. Khaddam-Aljameh, M. Dazzi, I. Giannopoulos, G. Karunaratne *et al.*, "Deep learning acceleration based on in-memory computing," *IBM Journal of Research and Development*, vol. 63, no. 6, pp. 7–1, 2019.
- [175] S. Woźniak, A. Pantazi, T. Bohnstingl, and E. Eleftheriou, "Deep learning incorporating biologically inspired neural dynamics and in-memory computing," *Nature Machine Intelligence*, vol. 2, no. 6, pp. 325–336, 2020.
- [176] A. Sebastian, M. Le Gallo, R. Khaddam-Aljameh, and E. Eleftheriou, "Memory devices and applications for in-memory computing," *Nature nanotechnology*, vol. 15, no. 7, pp. 529–544, 2020.
- [177] I. Fedorov, R. P. Adams, M. Mattina, and P. Whatmough, "Sparse: Sparse architecture search for cnns on resource-constrained microcontrollers," *Advances in Neural Information Processing Systems*, vol. 32, 2019.
- [178] C. Banbury, C. Zhou, I. Fedorov, R. Matas, U. Thakker, D. Gope, V. Janapa Reddi, M. Mattina, and P. Whatmough, "Micronets: Neural network architectures for deploying tinyml applications on commodity microcontrollers," *Proceedings of Machine Learning and Systems*, vol. 3, pp. 517–532, 2021.
- [179] B. Sudharsan, J. G. Breslin, and M. I. Ali, "ML-mcu: A framework to train ml classifiers on mcu-based iot edge devices," *IEEE Internet of Things Journal*, 2021.
- [180] R. Dai, L. Shen, F. He, X. Tian, and D. Tao, "Displf: Towards communication-efficient personalized federated learning via decentralized sparse training," *arXiv preprint arXiv:2206.00187*, 2022.
- [181] A. Li, J. Sun, P. Li, Y. Pu, H. Li, and Y. Chen, "Hermes: an efficient federated learning framework for heterogeneous mobile clients," in *Proceedings of the 27th Annual International Conference on Mobile Computing and Networking*, 2021, pp. 420–437.

- [182] E. Diao, J. Ding, and V. Tarokh, "Heterofl: Computation and communication efficient federated learning for heterogeneous clients," in *International Conference on Learning Representations*, 2020.
- [183] J. Zhang, S. Guo, X. Ma, H. Wang, W. Xu, and F. Wu, "Parameterized knowledge transfer for personalized federated learning," *Advances in Neural Information Processing Systems*, vol. 34, pp. 10 092–10 104, 2021.
- [184] N. Volgushev, M. Schwarzkopf, B. Getchell, M. Varia, A. Lapets, and A. Bestavros, "Conclave: secure multi-party computation on big data," in *Proceedings of the Fourteenth EuroSys Conference 2019*, 2019, pp. 1–18.